

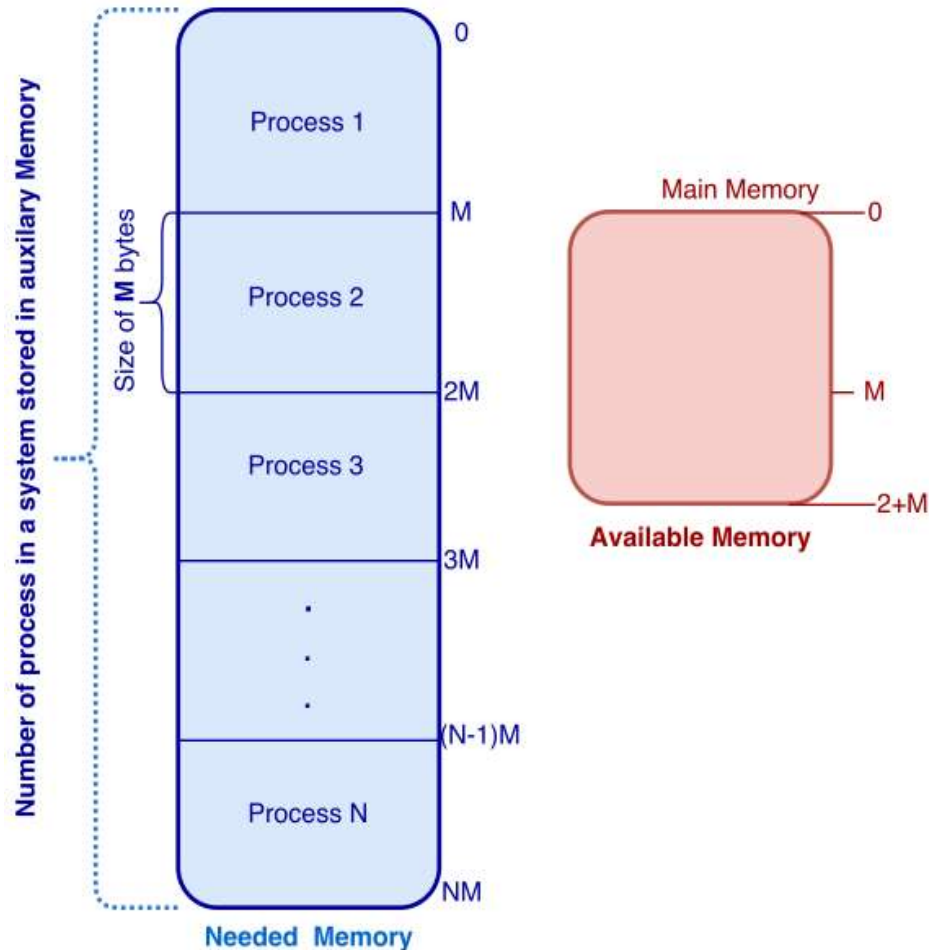
Understanding Memory Management

L J Gokul Vasan

Contents

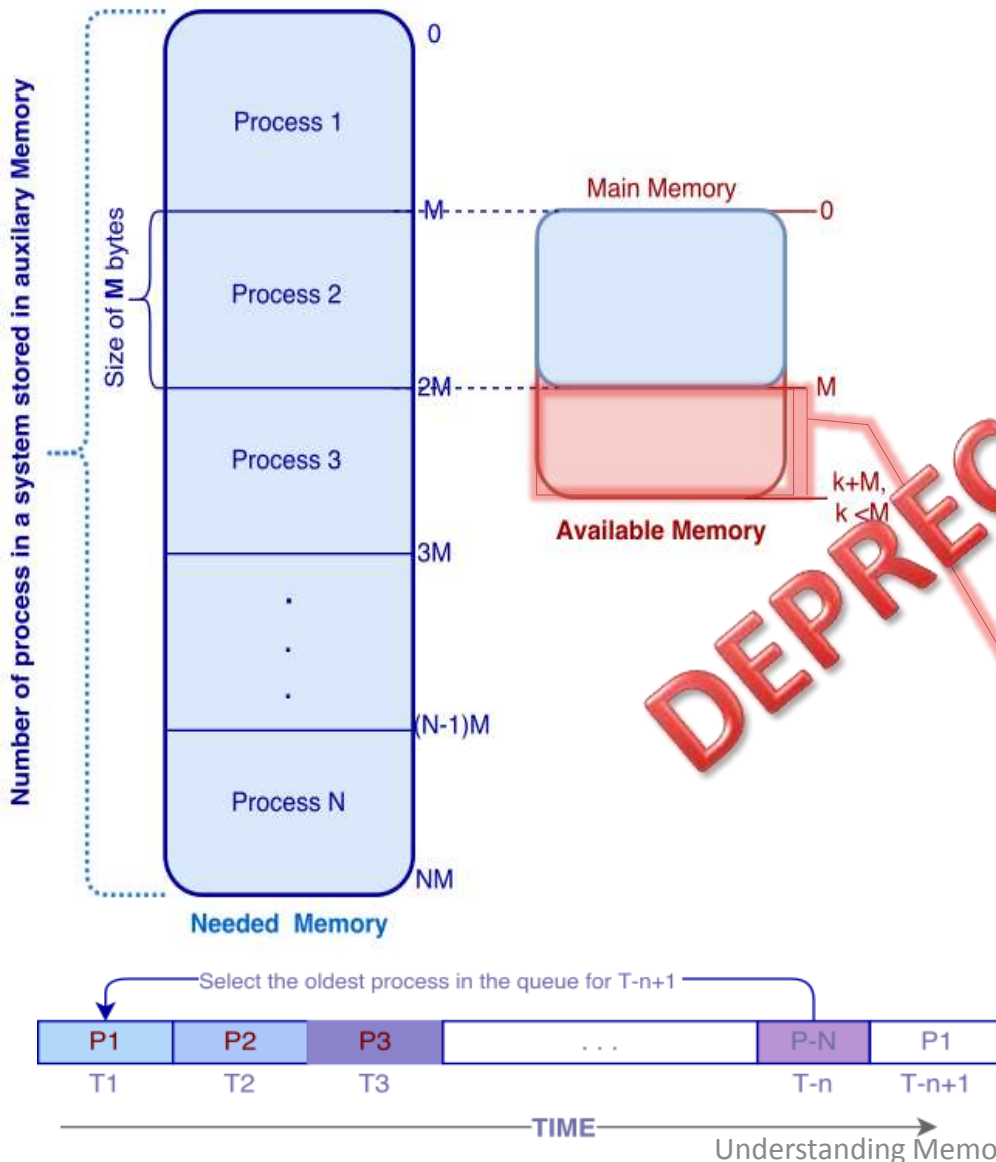
- Basics
 - Swapping
 - Paging
 - Virtual Memory
- Hardware in memory management
 - TLB
 - Translation of address
- Theory on memory management
 - Locality of reference
 - Thrashing Problem
 - Working set
- Software in Memory management
 - Replacement policy
 - Placement policy
 - Scan rate policy
 - Fetch Policy

Process and Memory



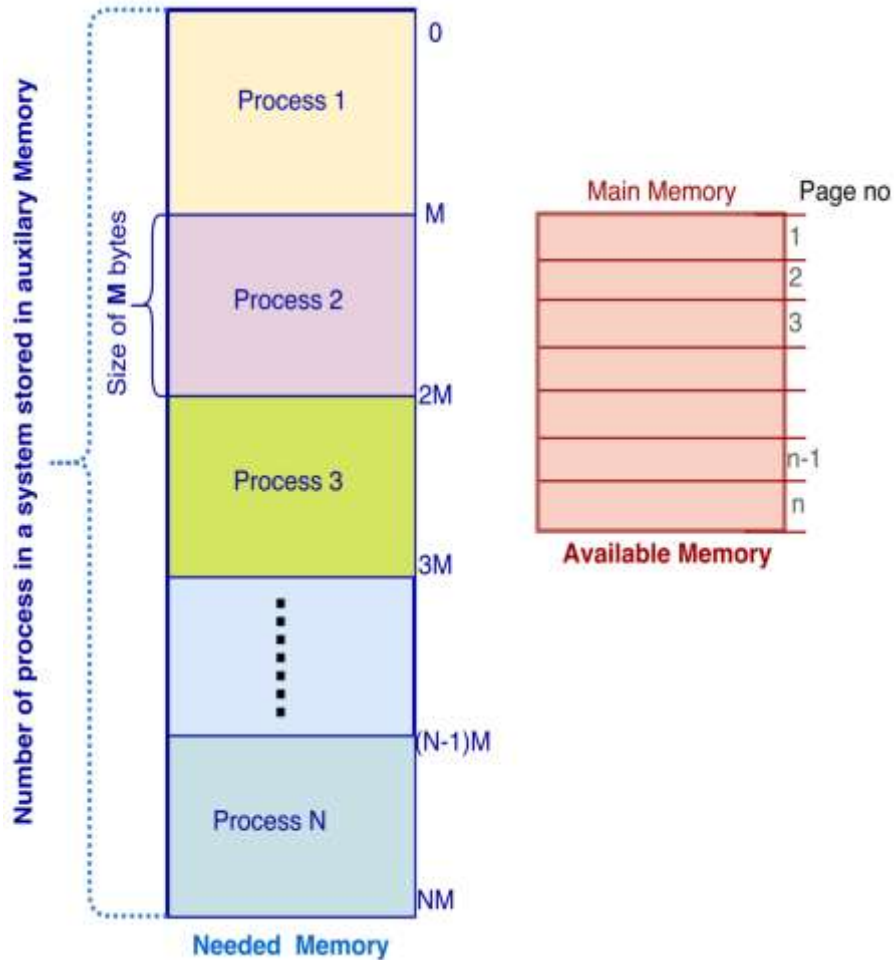
- In reality, main memory is smaller than required memory.

Swapping



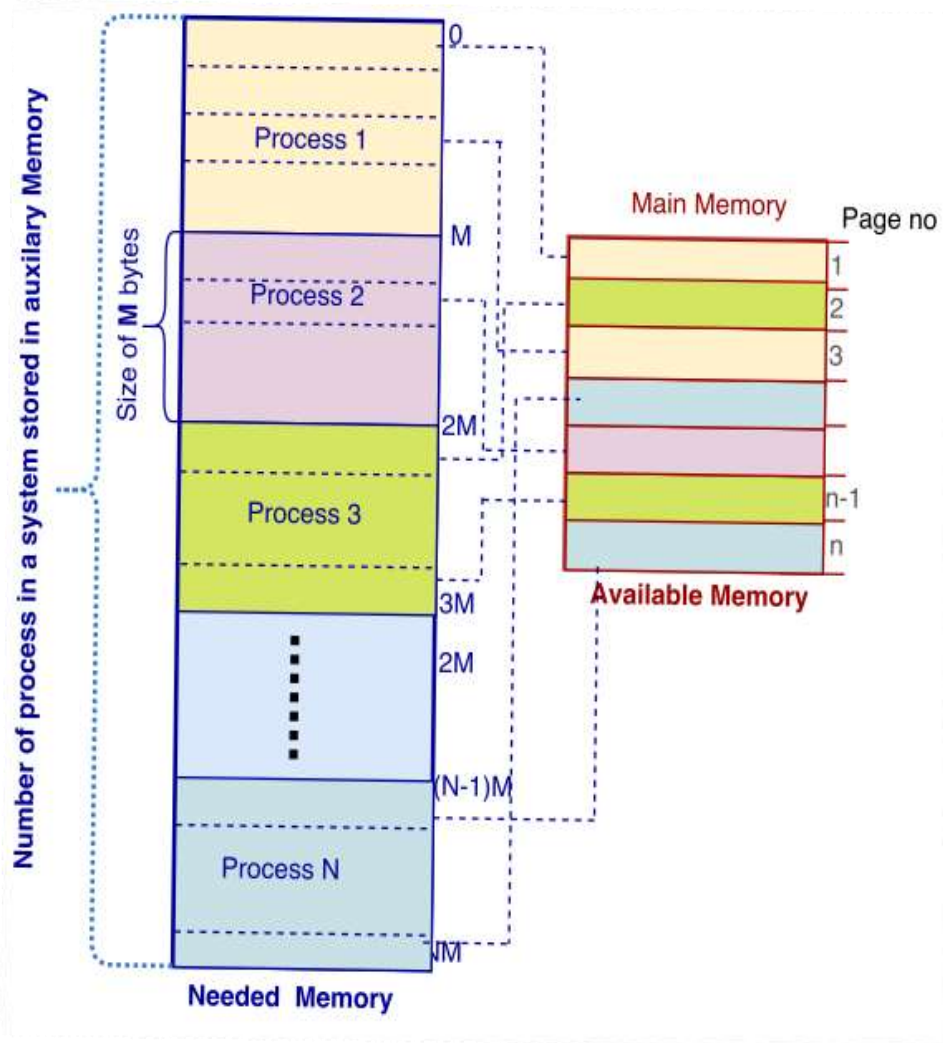
- A process in main memory is **swapped** on time sharing or priority basis.
- **Swapping**: The whole process is brought in and out of main memory.
- **Predicament**:
 - External Fragmentation.
 - larger process can exist than main memory.
 - Time Consuming

Paging



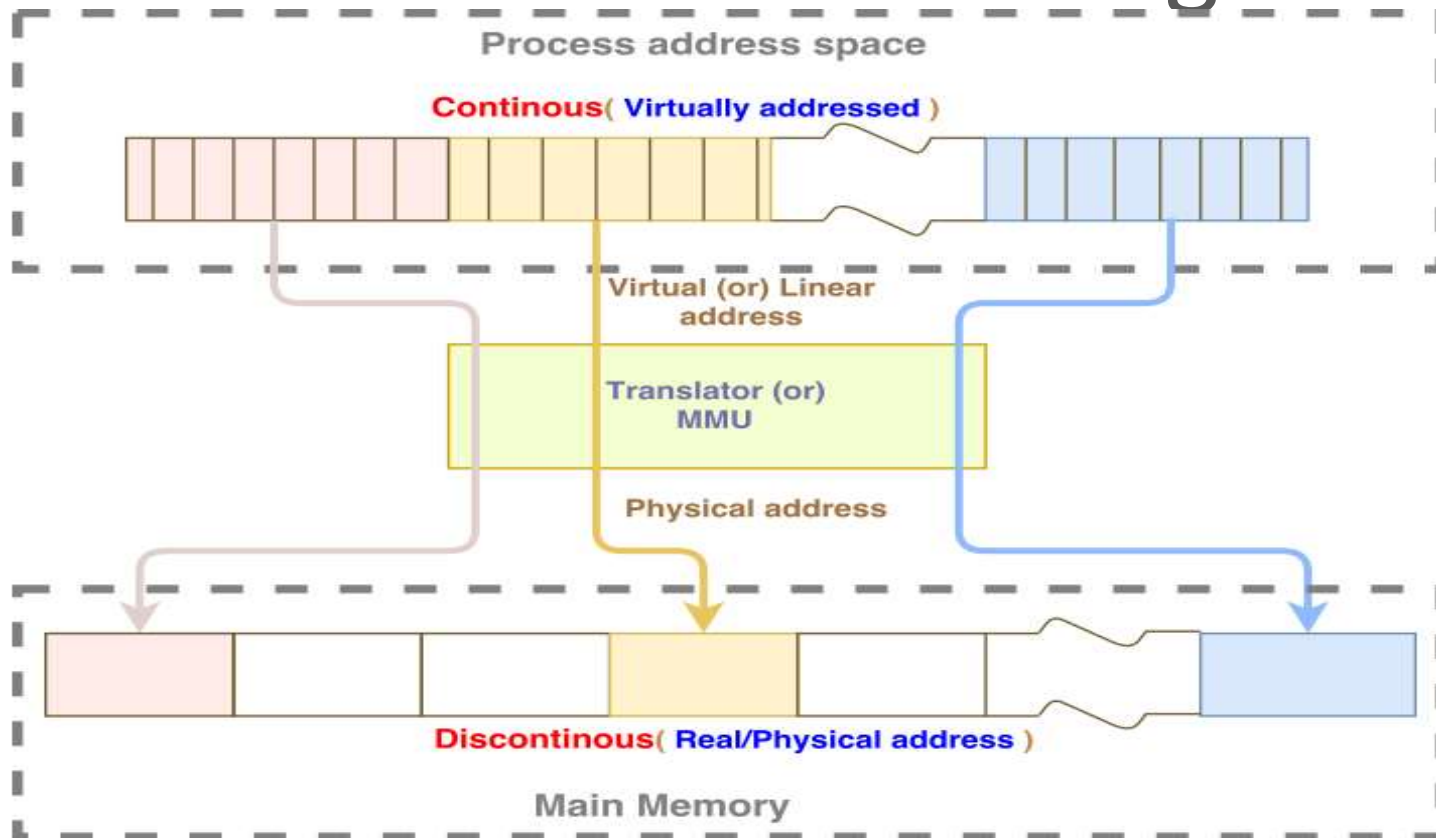
- Main memory is divided into equal chunks called **pages**.
- Only a piece of the process is loaded into the main memory.
- **Paging**: On Exhaustion of main memory only a **page or a set of pages** are evicted from main memory.

Paging



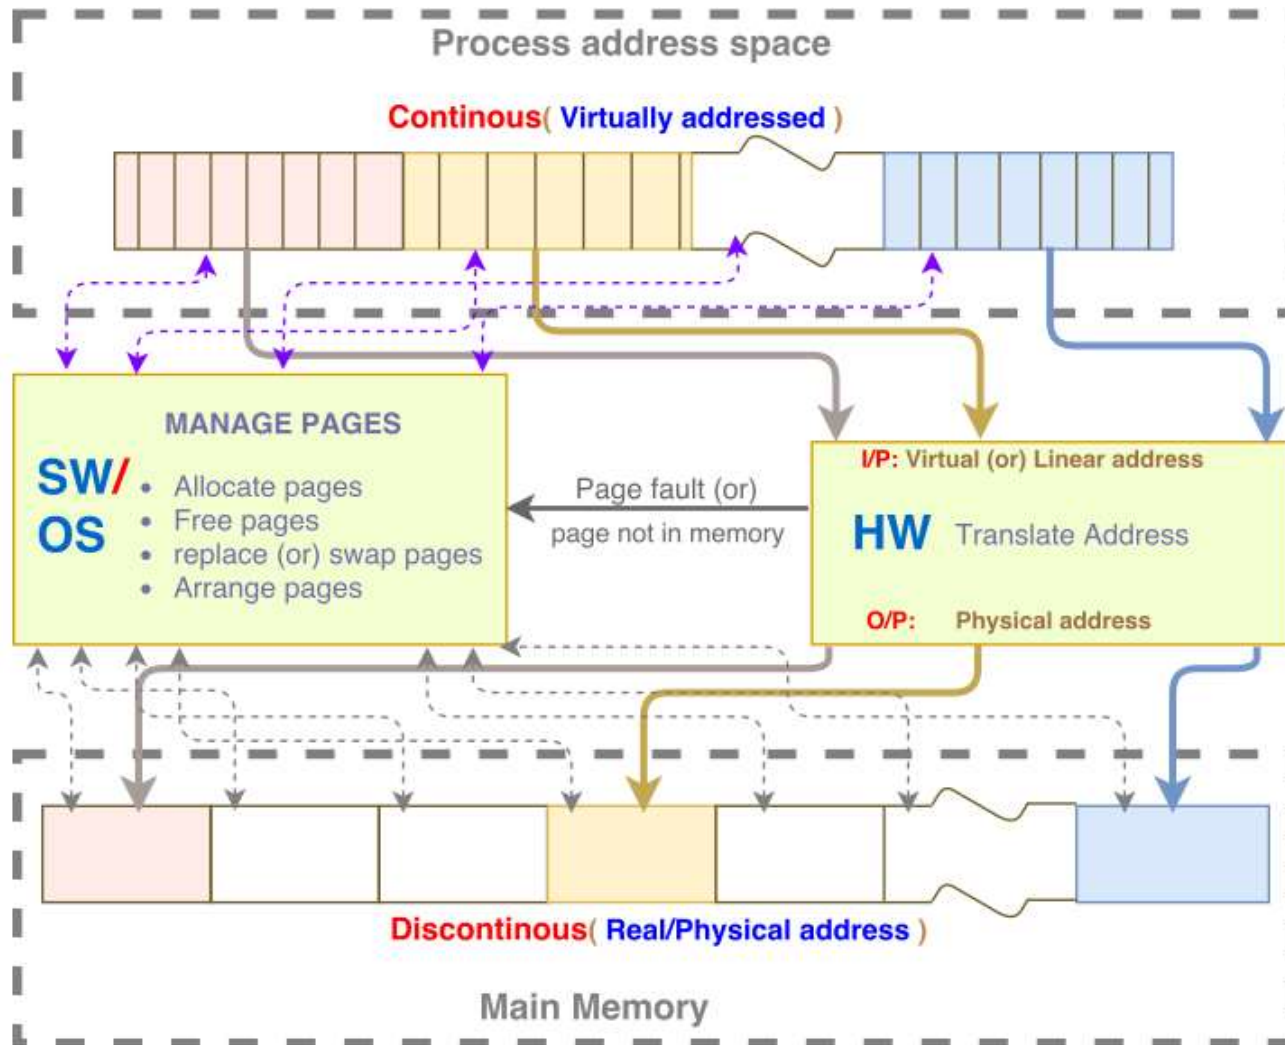
- **Advantages:**
 - Selective loading of a process.
 - process could load on **discontinuous pages**, avoiding external fragmentation.
 - Process can be larger than main memory.
- **Problem:**
 - A Process cannot be addressed continuously.
- **Solution:**
 - Virtual addressing.

Virtual Addressing



- Now, there are **2 addresses**:
 - Virtual (or) Linear address (**simulated to provide continuous perspective**).
 - Physical address (**represents real memory**).
- Process is virtually addressed.
- Translator named **Memory management unit(MMU)** converts Virtual to Physical.

Role of OS and HW in Memory Management

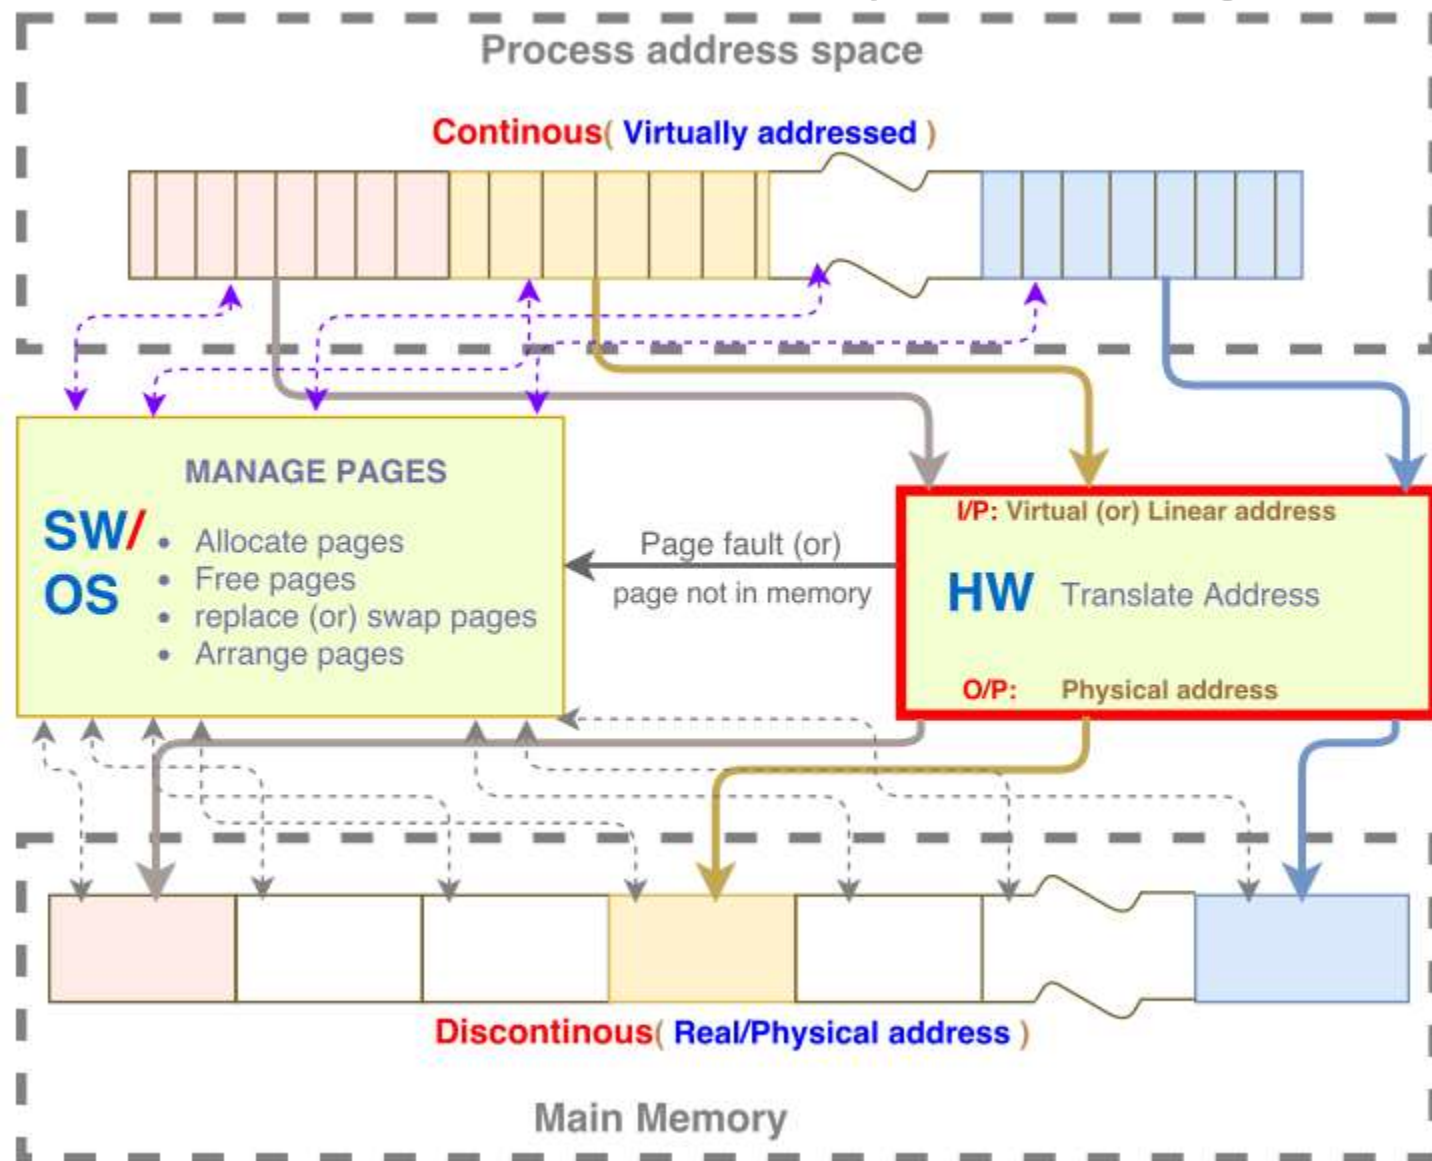


Hardware in Memory management


Address Translation

Virtual -> Physical

Hardware in Memory management

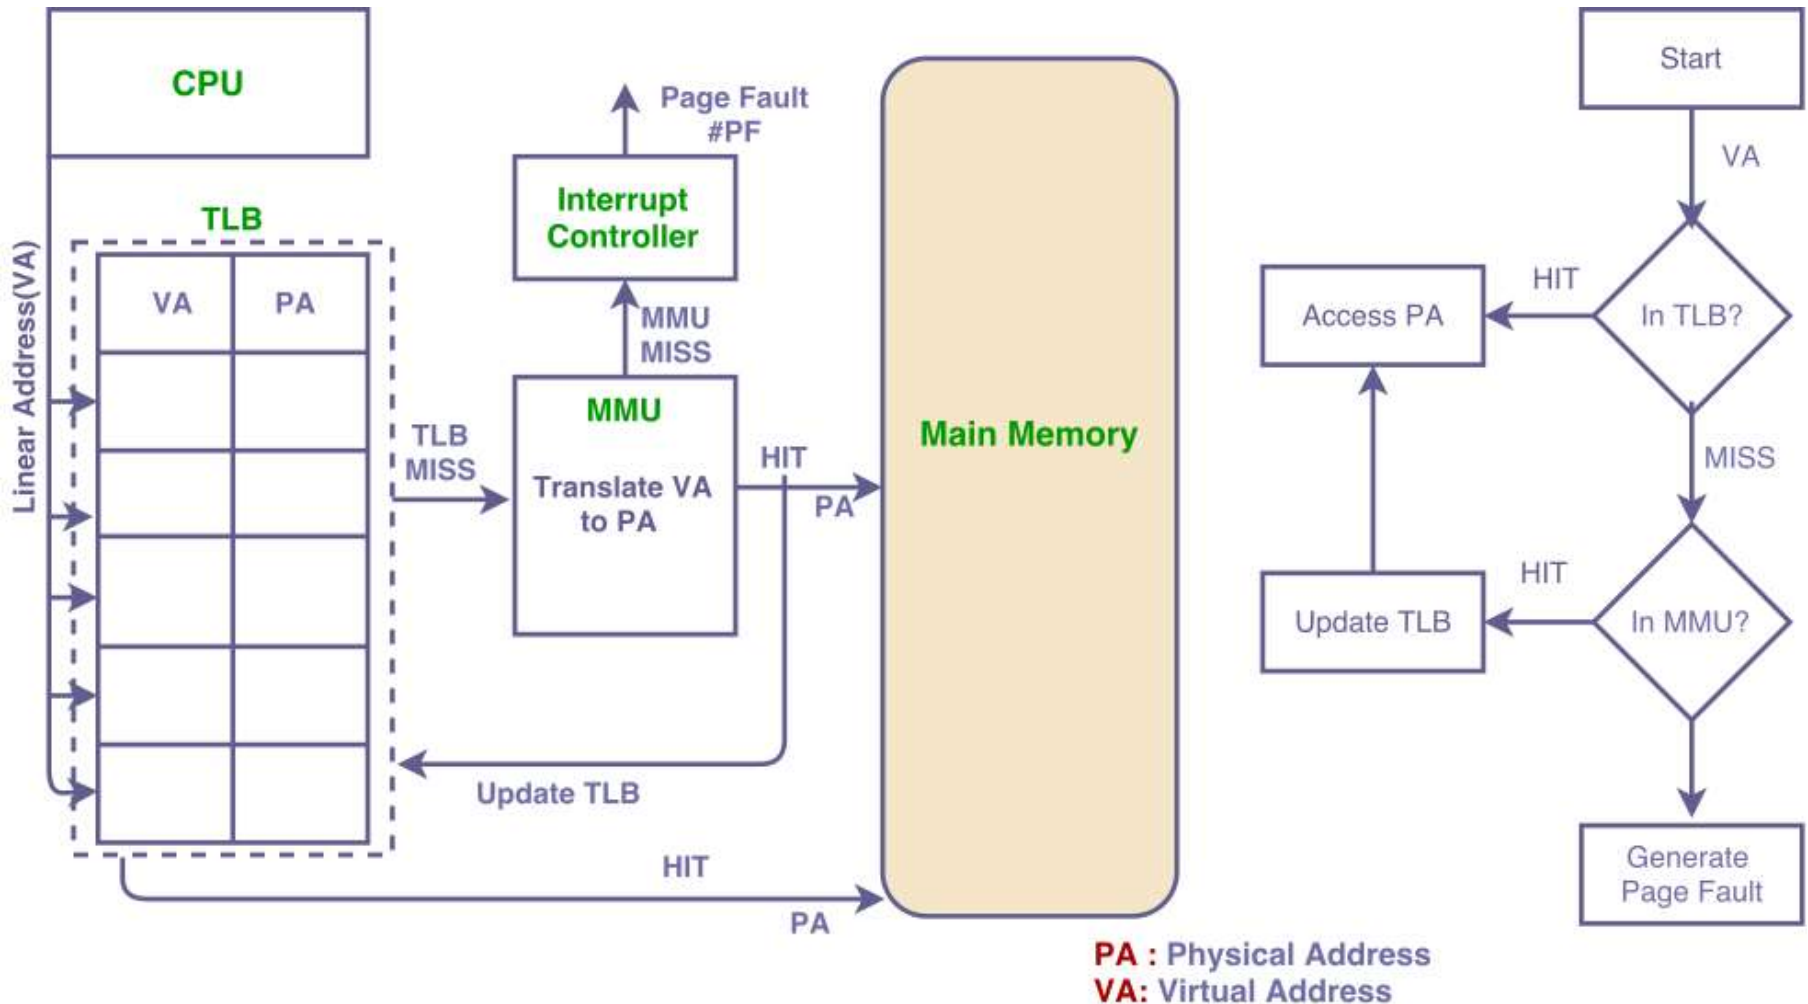


MMU position in HW




- Placed between CPU and memory controller.
- **TLB**(**Translation Lookaside Buffer**) would hold the last recent translated page addresses.
 - Reduces computation time of MMU translation.
- If not present in TLB, MMU translates and also updates TLB.

MMU and TLB data flow(Simplified)



Virtual Address

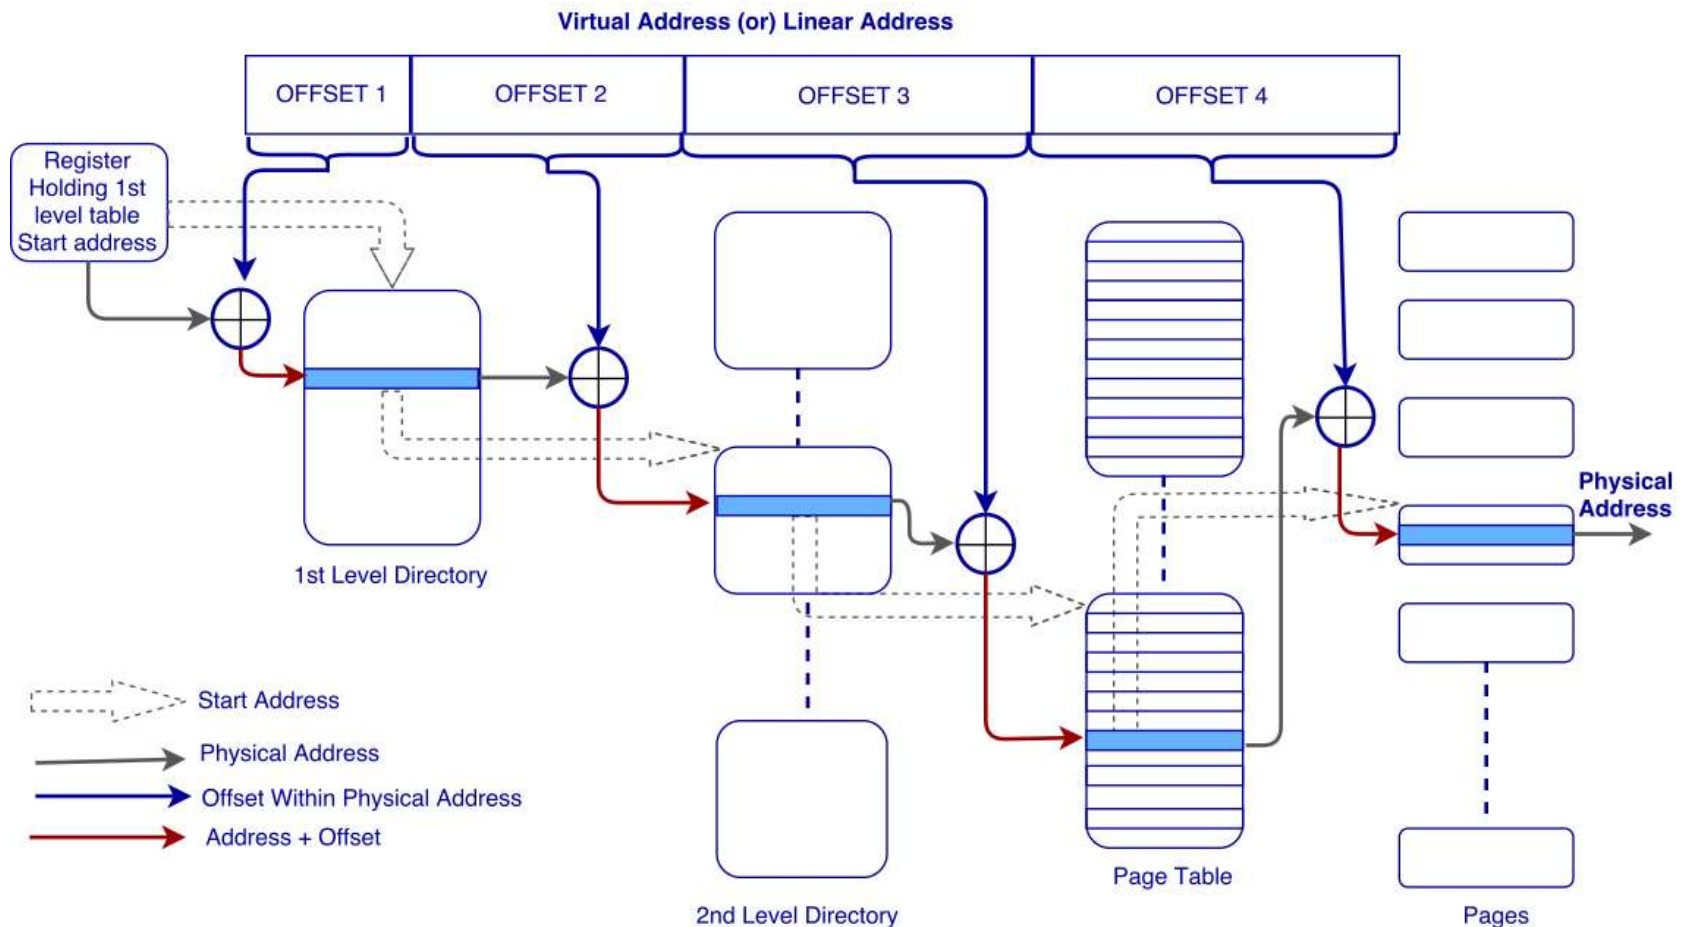
- Virtual address is a sequence of offsets.



- Offsets:** Complete address is divided into **smaller contiguous bit chunks**.
- Each bit chunk (Offset) holds a offset or index of a **table**.
 - Table will hold the starting address of the next table.
 - This address is added with the next bit chunk to derive the right location in next table or page.
- Tables are **pages of memory**, however, rather holding the data it is a list holding address of next table or page frame.
- On schedule of a processes, the OS loads the process's **Directory start address (offset 1) into the corresponding memory management register**.

MMU: Translation Procedure(Generalised)

- On TLB miss MMU does translation




Why the concept of paging, TLB and caches work?

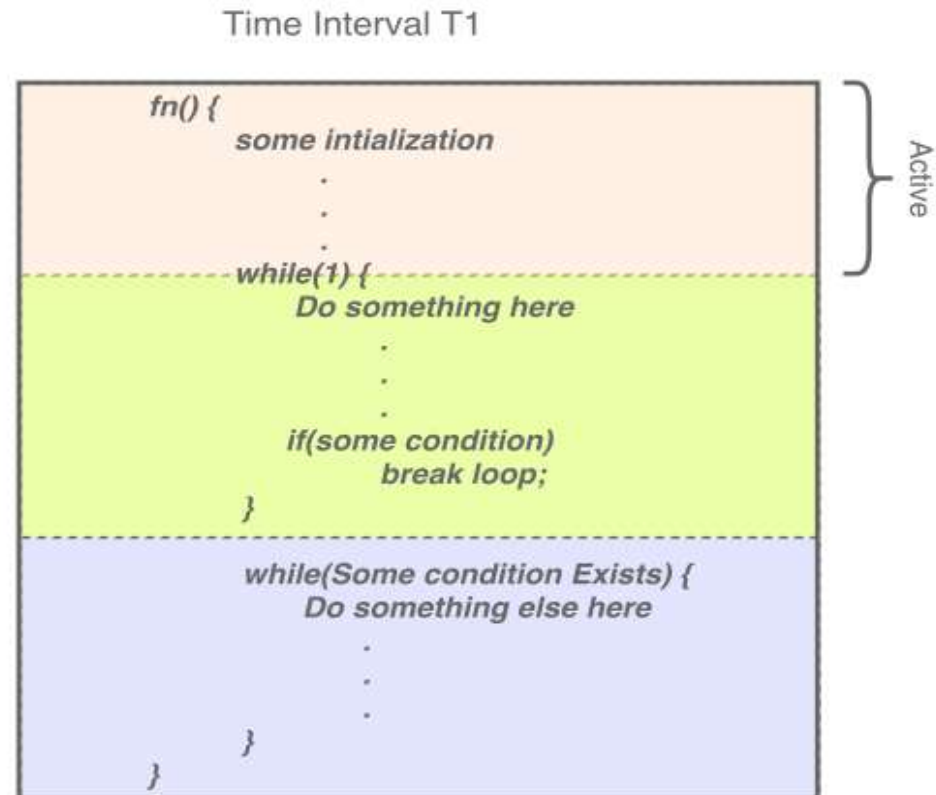
Locality of reference

Locality of reference(Exemplified)

- Time interval **T1**




Active Pages at
given time interval

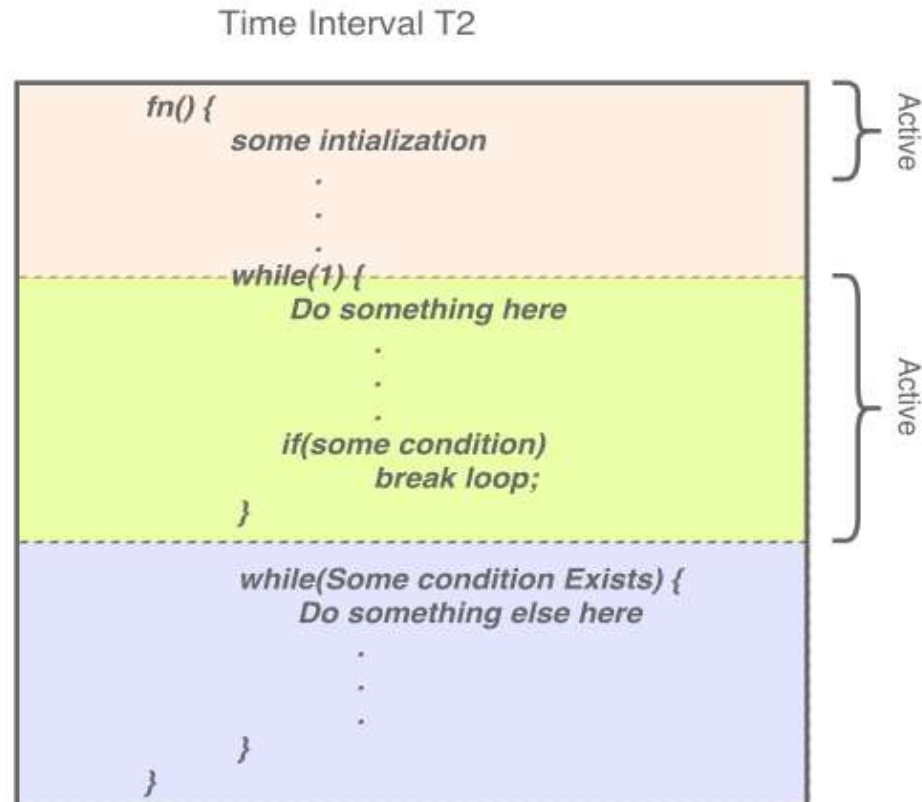


Locality of reference(Exemplified)

- Time interval **T2**




Active Pages at
given time interval

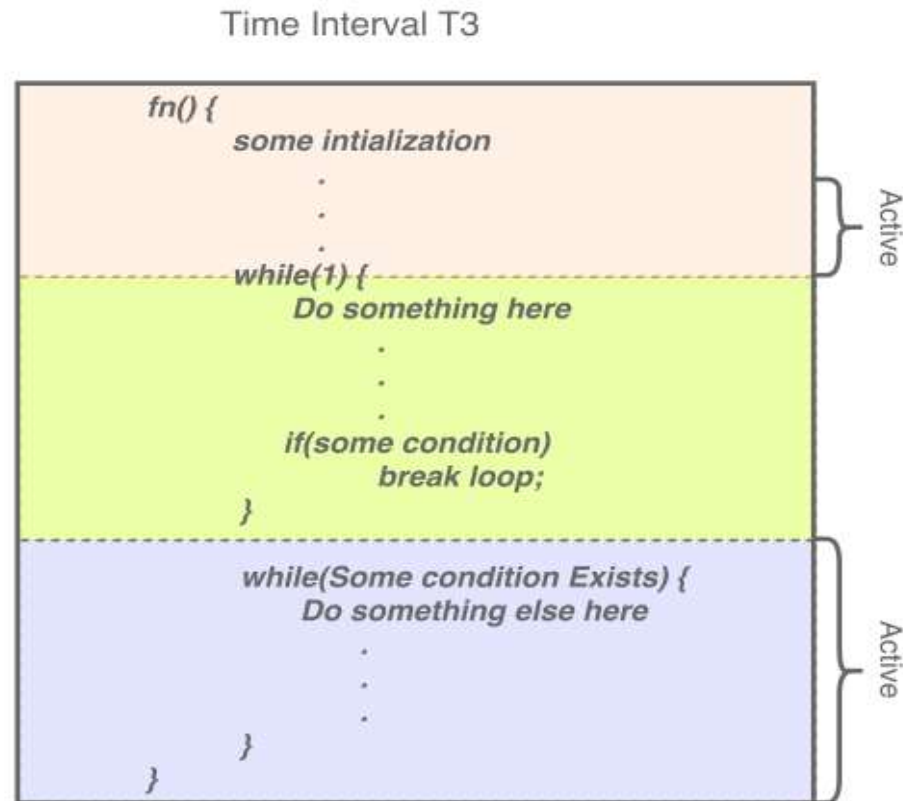


Locality of reference(Exemplified)

- Time interval **T3**



Active Pages at
given time interval



Locality of reference(In words)

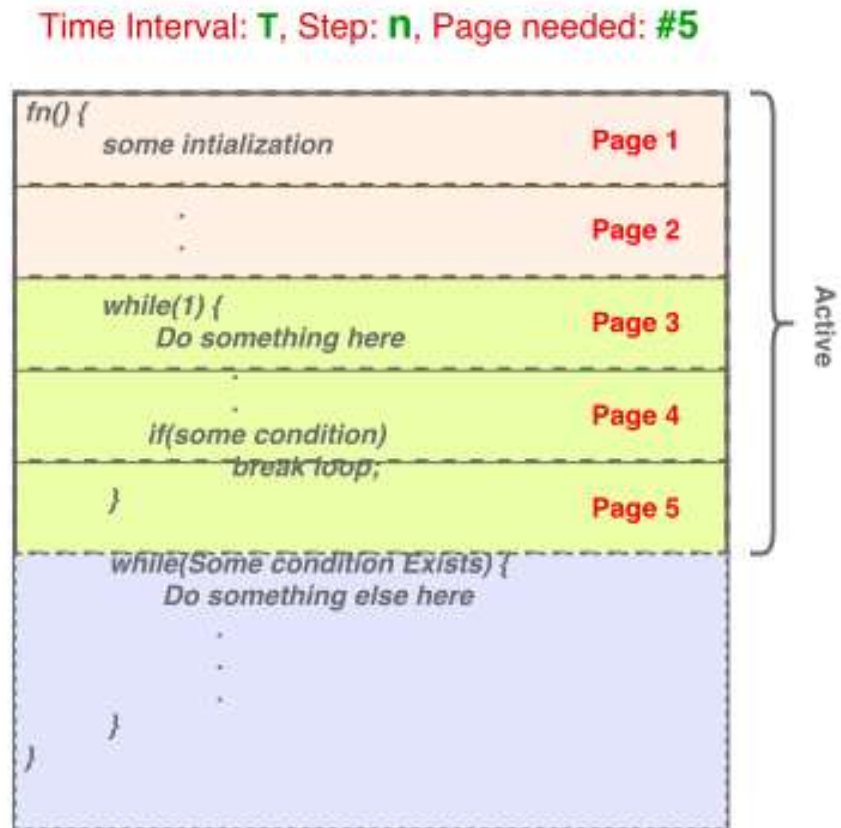
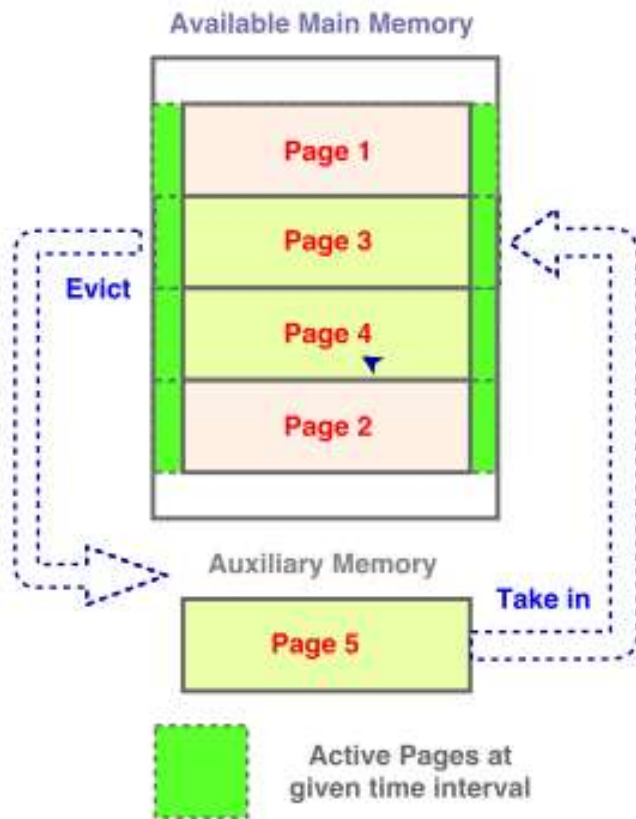
1. Its the tendency of a programme to **cluster references of pages to small set** of their pages for extended intervals.
2. There exists a **strong Relation between the near future and near past cluster** of reference pages, i.e., The set tends to overlap.
3. There exists a **feeble or nearly no relation between distant future and distant past** references of pages.
4. Pages are accessed in random exhibiting a **stochastic behaviour**.
5. The cluster references tends to slowly move away from one active set to another, i.e., They exhibit a **quasi stationary behaviour**, manifesting a **time series model**.

Thrashing problem

- A process spends significant computation on paging rather than on its real computation.

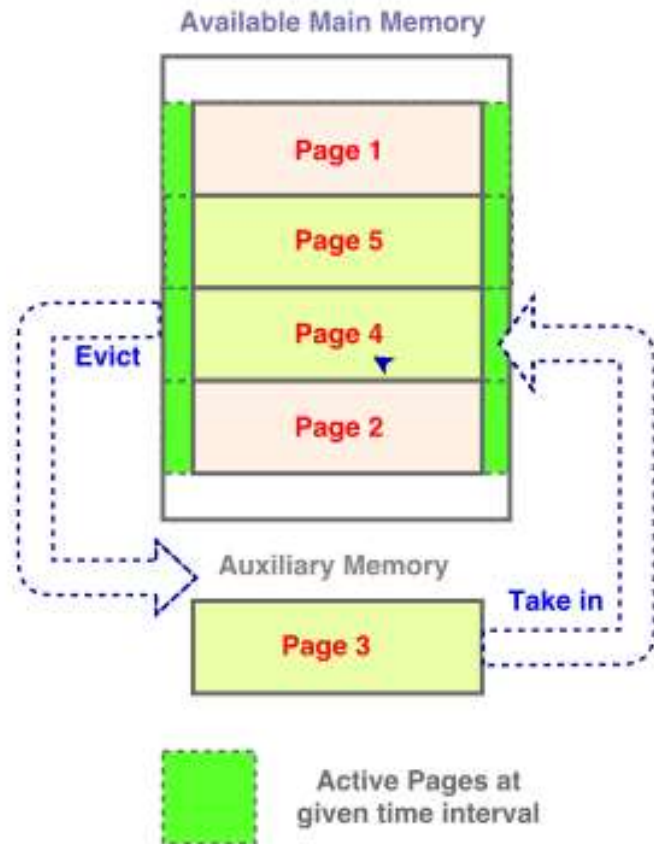
Thrashing(Exemplified)

- Required page count at this instance is **5** , but available pages is **4**.

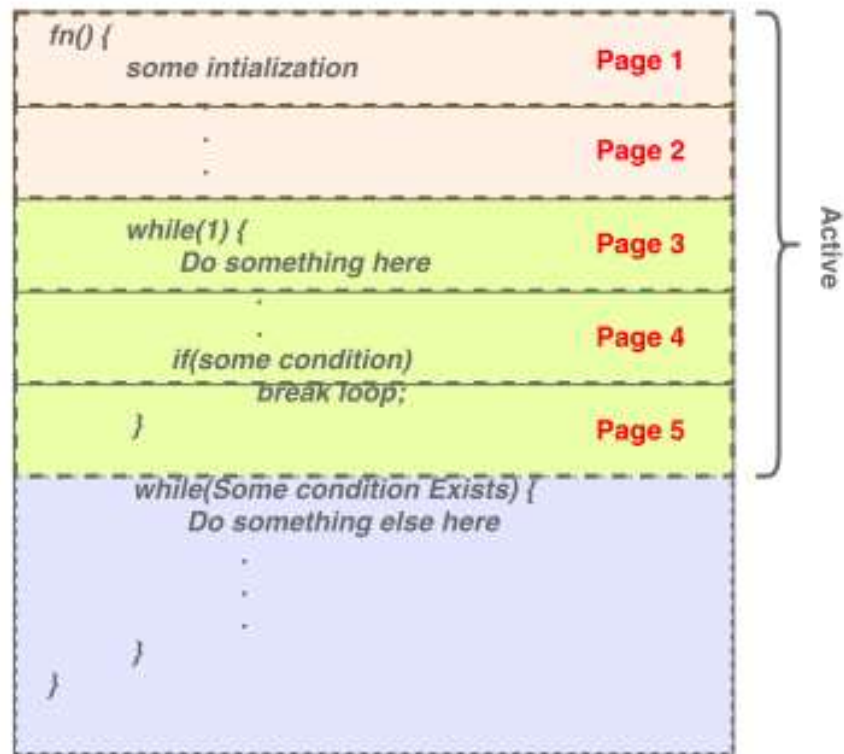


Thrashing(Exemplified)

- Required page count at this instance is **5** , but available pages is **4**.

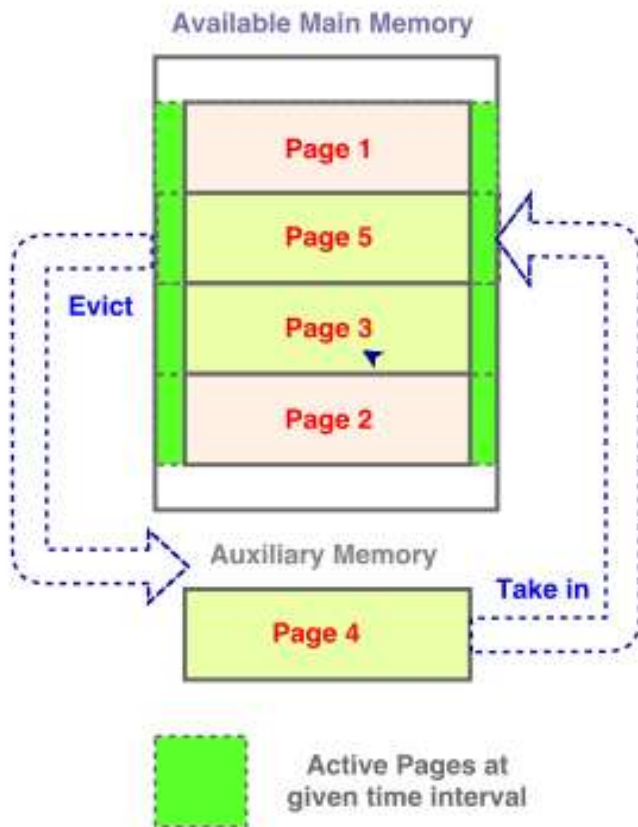


Time Interval: **T**, Step: **n+1**, Page needed: **#3**

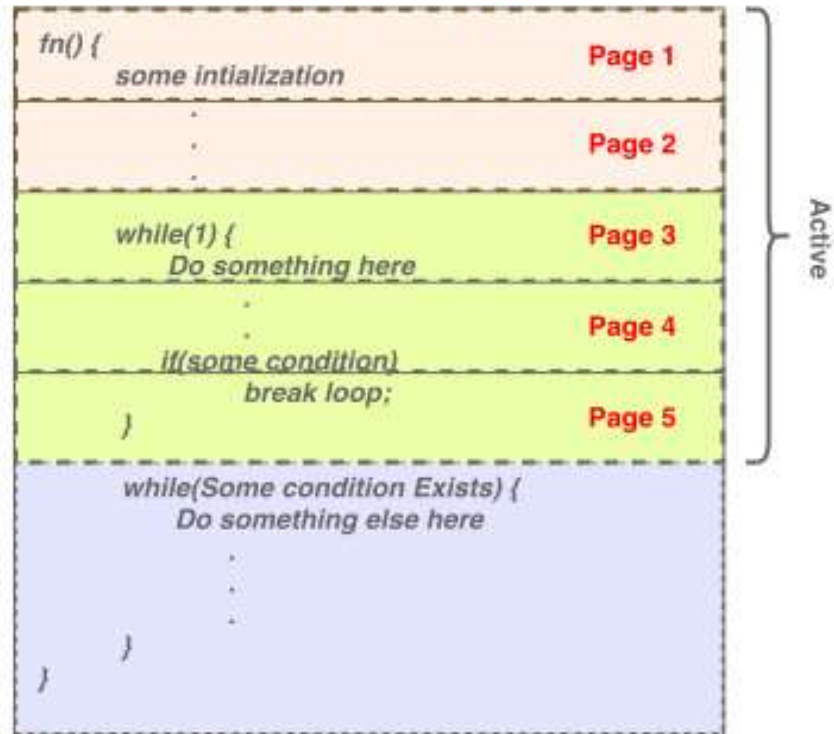


Thrashing(Exemplified)

- Required page count at this instance is **5** , but available pages is **4**.

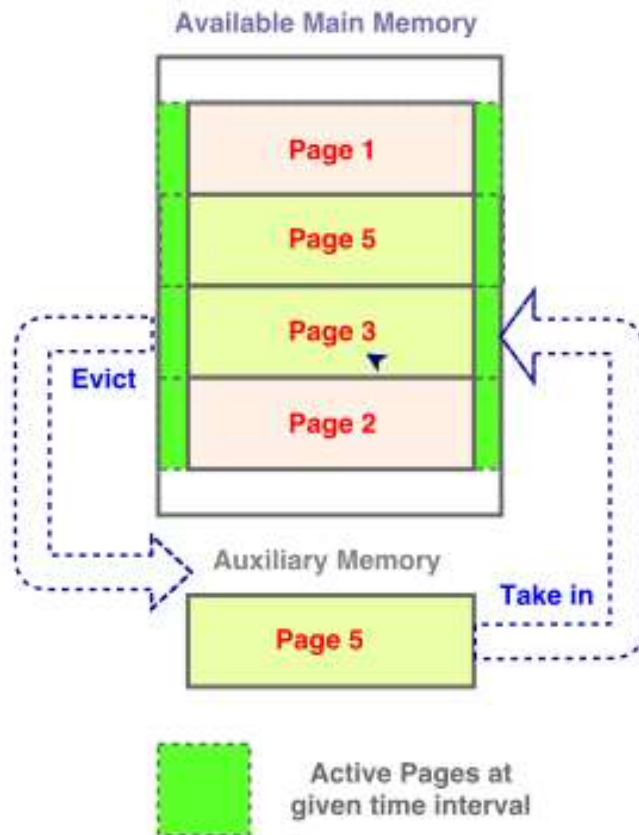


Time Interval: **T**, Step: **n+2**, Page needed: **#4**

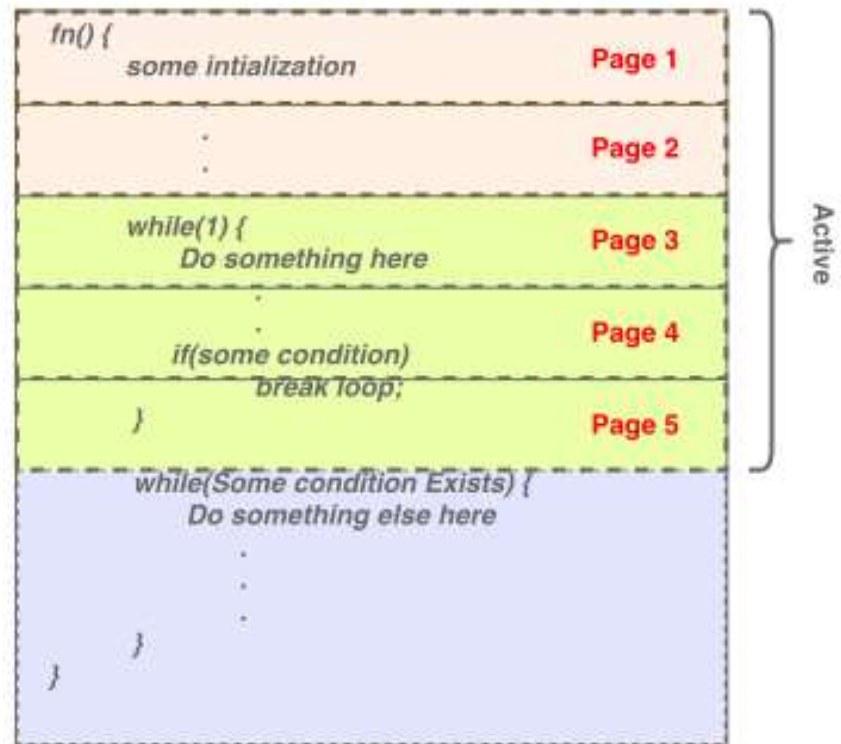


Thrashing(Exemplified)

- Required page count at this instance is **5** , but available pages is **4**.

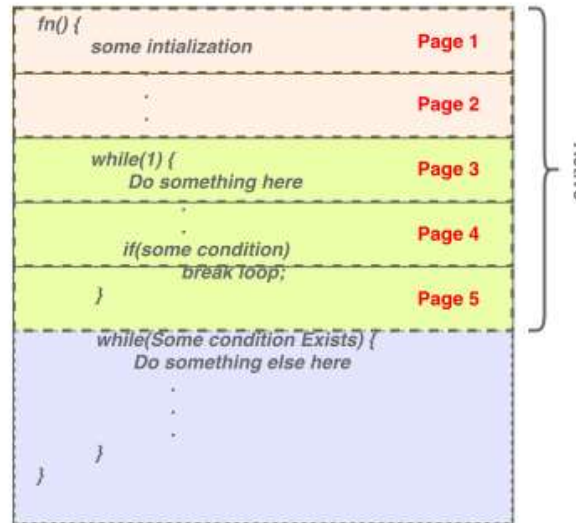
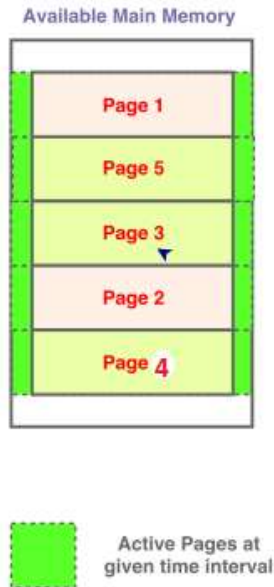


Time Interval: **T**, Step: **n+3**, Page needed: **#5**



Working set

- The working set model states that a process can be in main memory, **iff all of the pages that it is currently using can be in main memory.**



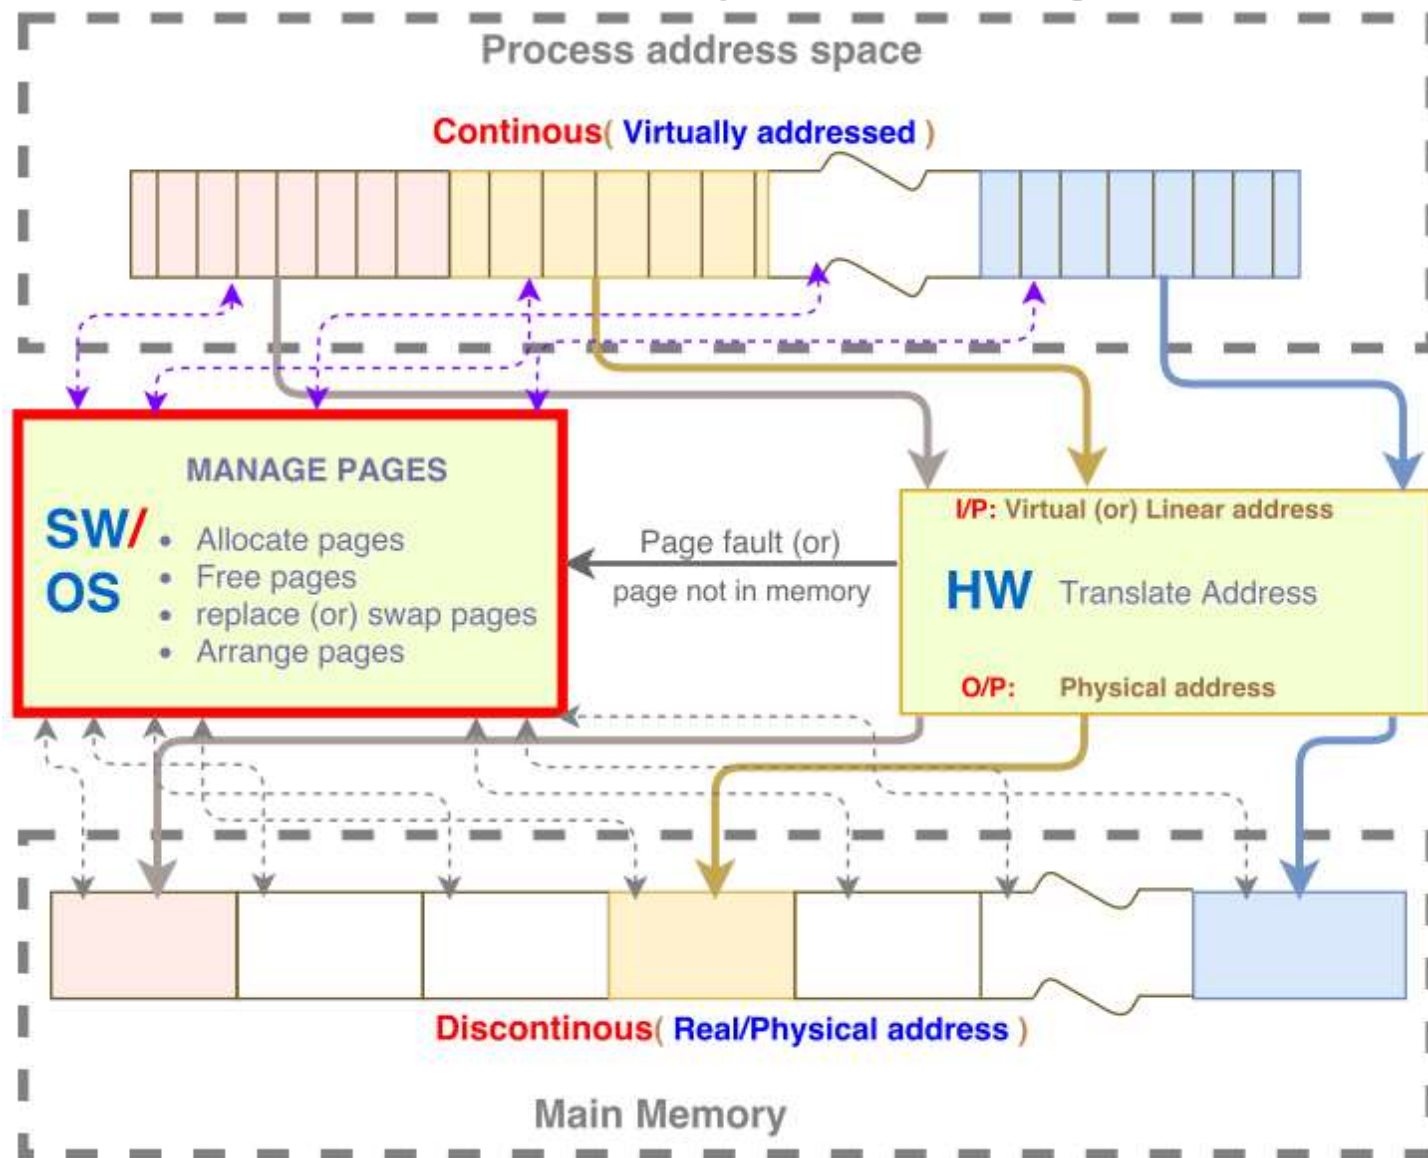
REQUIRED MEMORY: **5**
AVAILABLE MEMORY: **5**

Now, all pages needed by the programme is in memory, so no thrashing.

OS in Memory management

Policies (or) Rules


OS in Memory management



OS and MMU


- Hardware facilitates the virtual memory management, however OS determines on **optimal usage** of this ability.
- **Optimal usage:** Employ appropriate decision policies on
 - Which page to choose on allocation request (**Placement policy**).
 - Which page to evict on memory exhaustion(**Replacement policy**).
 - When to take in a page(**Fetch policy**).
 - Sort the pages for replacement(**Scan rate policy**).
- To execute the policies, OS needs **references** to physical pages.

Pages and Page Frame




- Divided chunks in main memory is termed as **Page frame**.
- Each page frame has **one-to-one** record associated in OS termed as **Page**.
- Pages enable OS to entail policies on page management.

Pages, Page frame and Zones




- Pages are further clustered into **zones**.
- zones are determined based on its location in physical memory and other criteria.
- Enables OS to apply different notions of allocation request.
- **Example:** Some **DMA devices are 16bit** can only access lower 65Kb, this would become a zone from which only 16bit DMA allocation request is fulfilled.

Pages and Active list: Ideal scenario



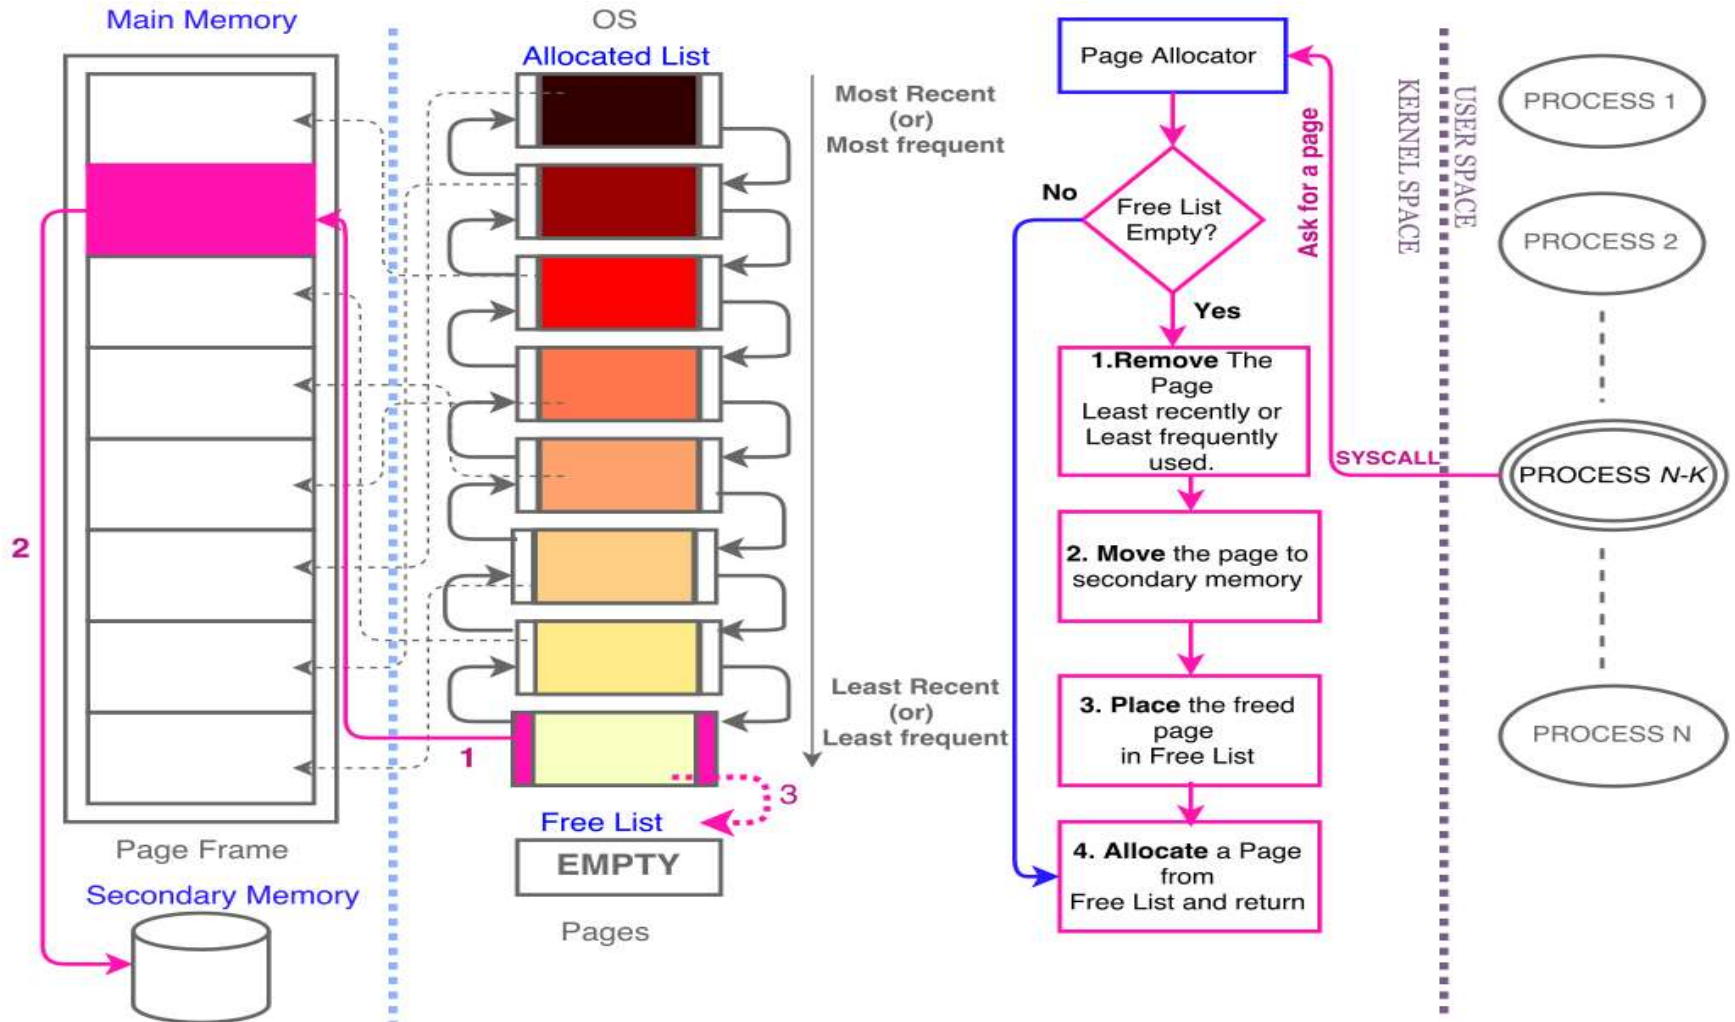
- Pages are kept in linked list.
- Initially, pages would be in free list.
- On request for memory, page is allocated from free list and moved to **allocated or active list**.
- quintessentially, pages in active list are sorted in descending order based on access or frequency.

Pages and Active list: Exhausted scenario



- Main memory could reach a point of exhaustion or near the point of exhaustion.
- After this point, if a process needs more memory, then system should not fail.
- **Replacement Policy:** OS should evict an **appropriate** page from main memory to auxiliary memory to pave way for new memory requirement.

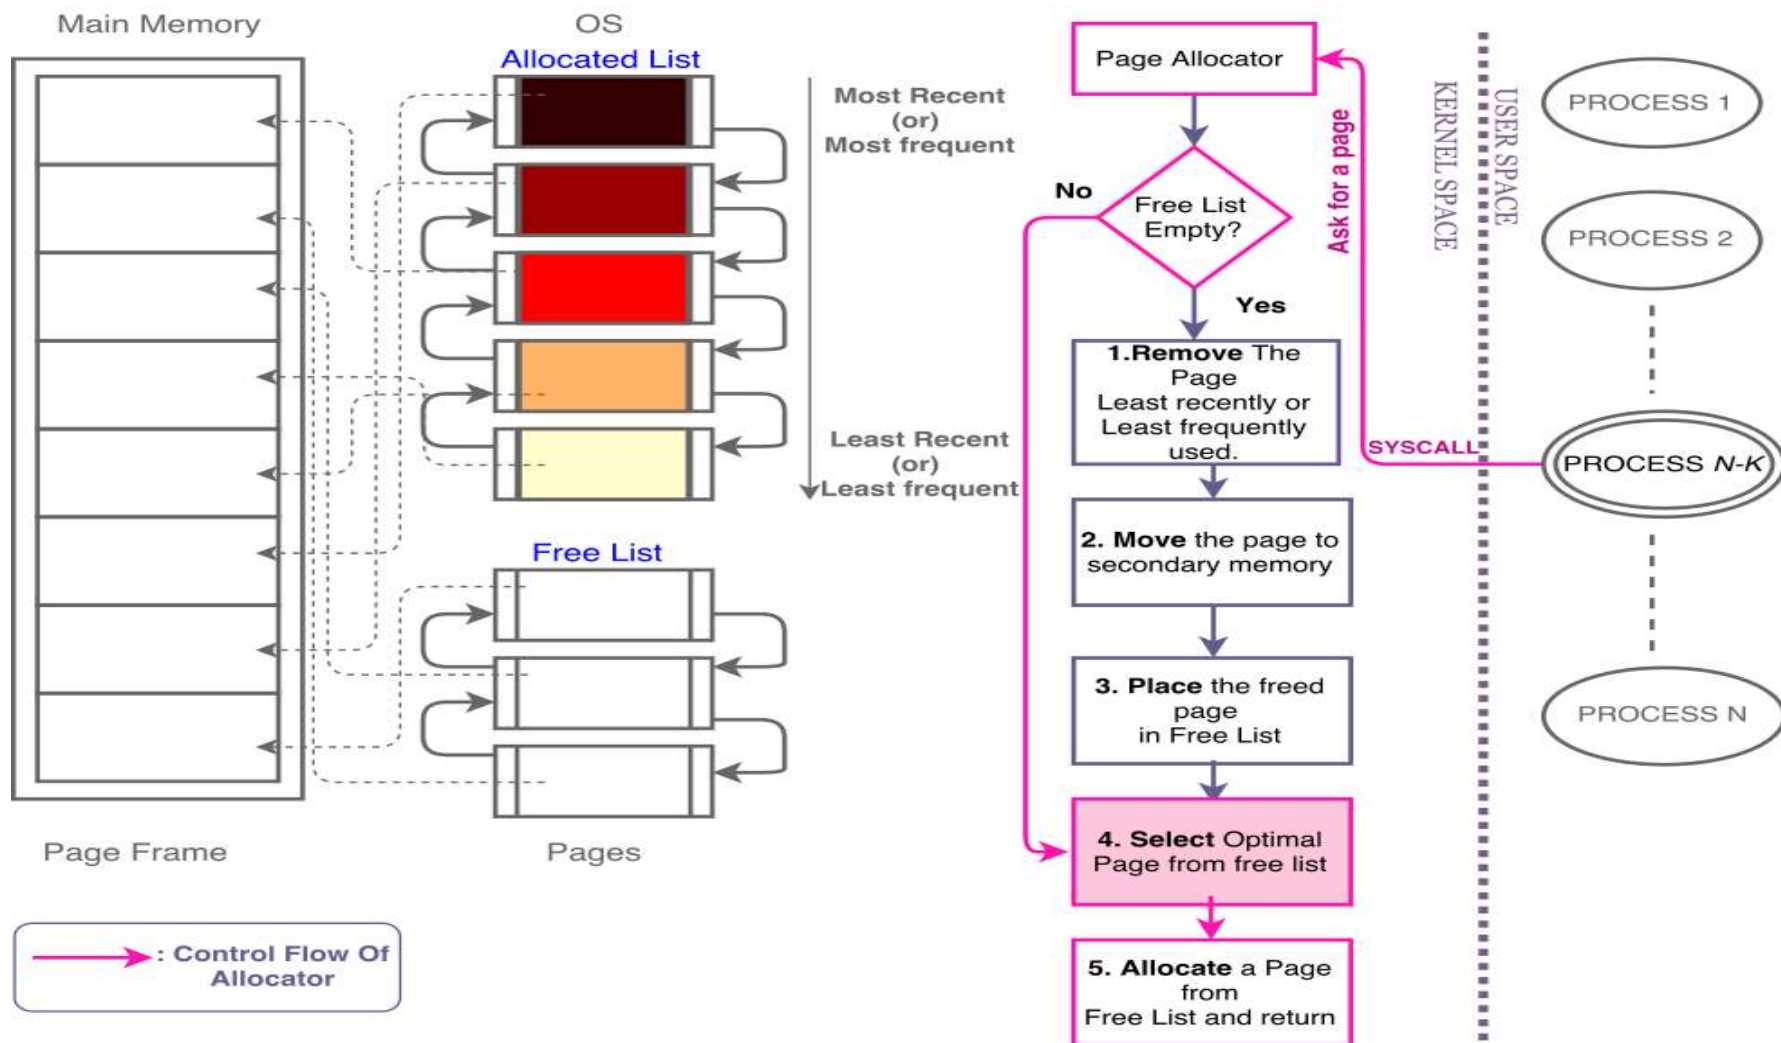
Replacement Policy(Simplified)



Placement Policy

- Tries to address following questions during allocation.
 - Are some physical pages **preferable** to others within a process?
 - Does allocating such a preferable page **increases performance** of the process?
- Addresses these query by selecting **optimal page** from the free list.


Placement Policy(Simplified)



Placement Policy: Optimal Page


- Pick a right page from free list such that:
 - It reduces **cache conflict or thrashing**. (or/and)
 - It is **near to the process's processor** (NUMA). (or/and)
 - It **saves energy**. (or/and)
 - Increases **parallelization**.

Active List: Reality




- In reality, pages in active list **will not be sorted**.
- In General, a system might have 1000s of active pages.
- **Sorting** the entire list periodically is **a CPU intensive computation**.
- Might reduce the system's performance considerably.

Scan rate Policy



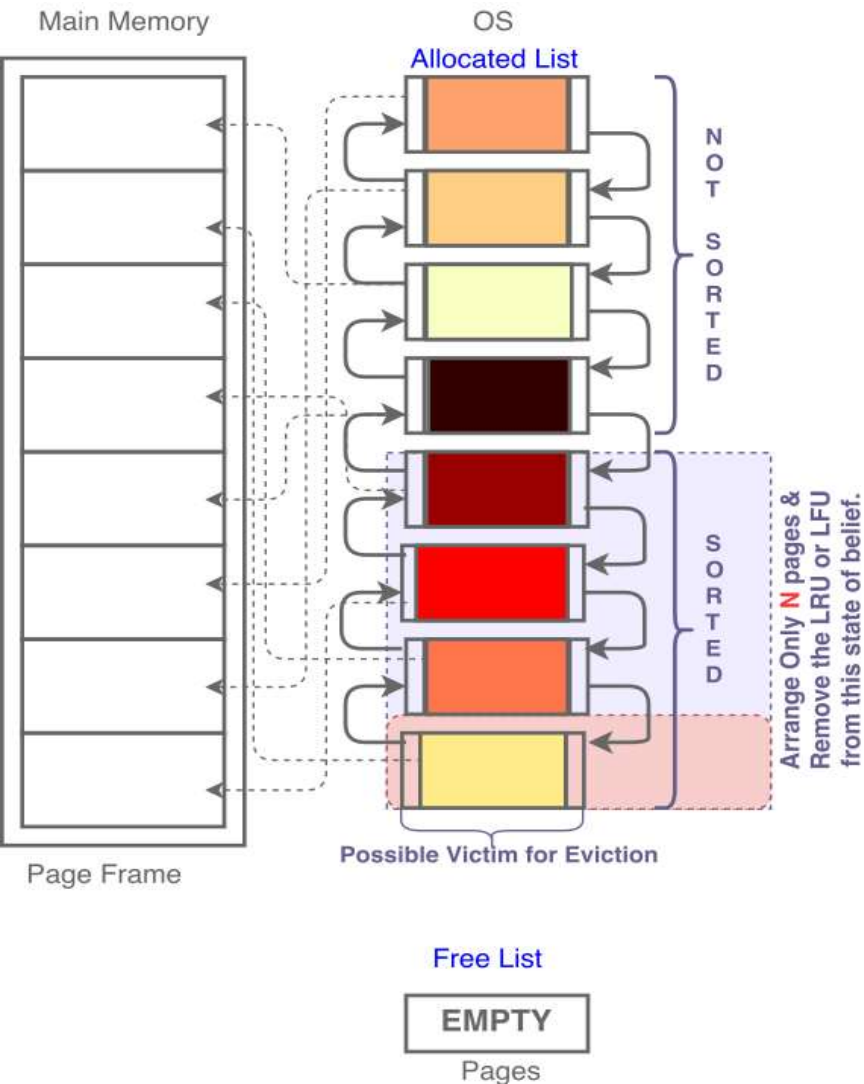
- **Scanning:** A process of arranging the pages in active list on accordance to replacement policy.
- **Scan rate policy:** Determines on scanning by deciding on 2 parameters
 - How much to scan.
 - When to scan.
- In general, **scanning and replacement are done jointly.**

Scan rate Policy: When to Scan



- Determines when to **start and stop scanning**.
- Determines **different scan rate policy to apply based on watermark**.
- **Watermark**: Zones have way points based on their distance from exhaustion.
- In general, watermark are **static**, determined by userspace variable (or) computed offline.
- Watermark also **determines how much to scan**.

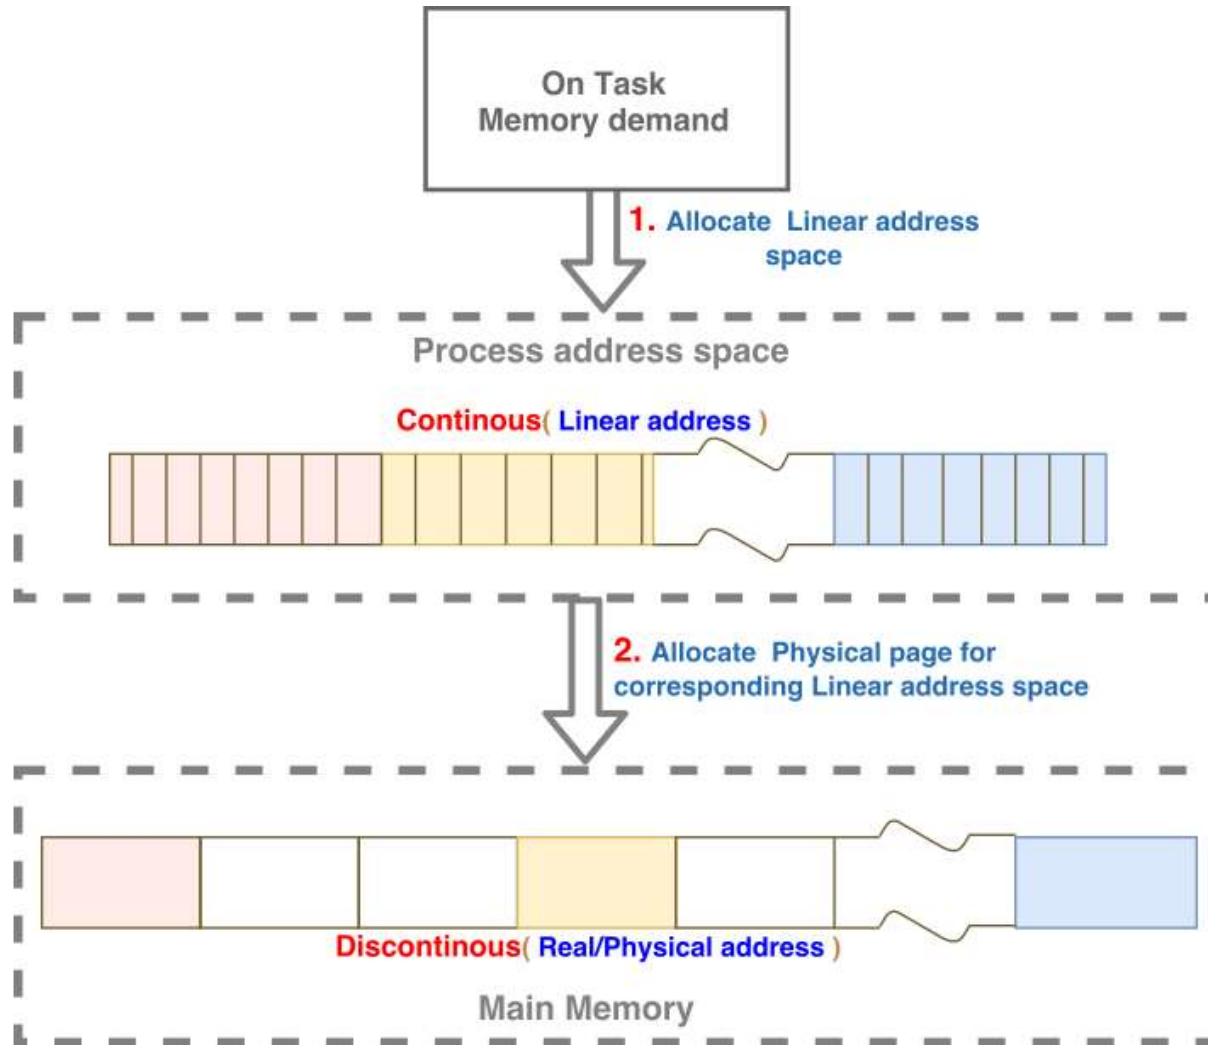
Scan rate Policy: How much to scan



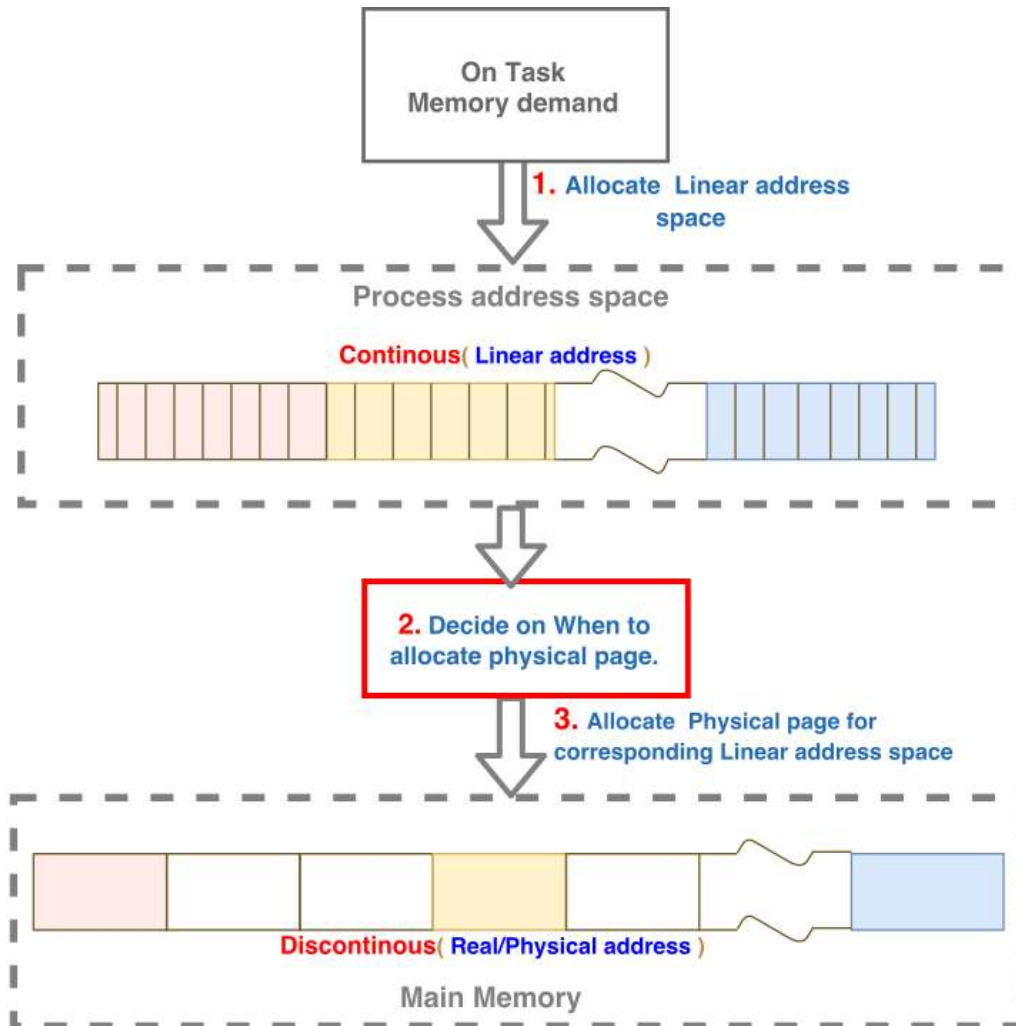
- Determines the measure on the range to scan within active list.
- Directly impacts the performance of the system.
- In general, nearer to exhaustion more the scanning, as more pages needs eviction.

Address and Memory in OS

OS divides **address and memory** into **two** different entities.



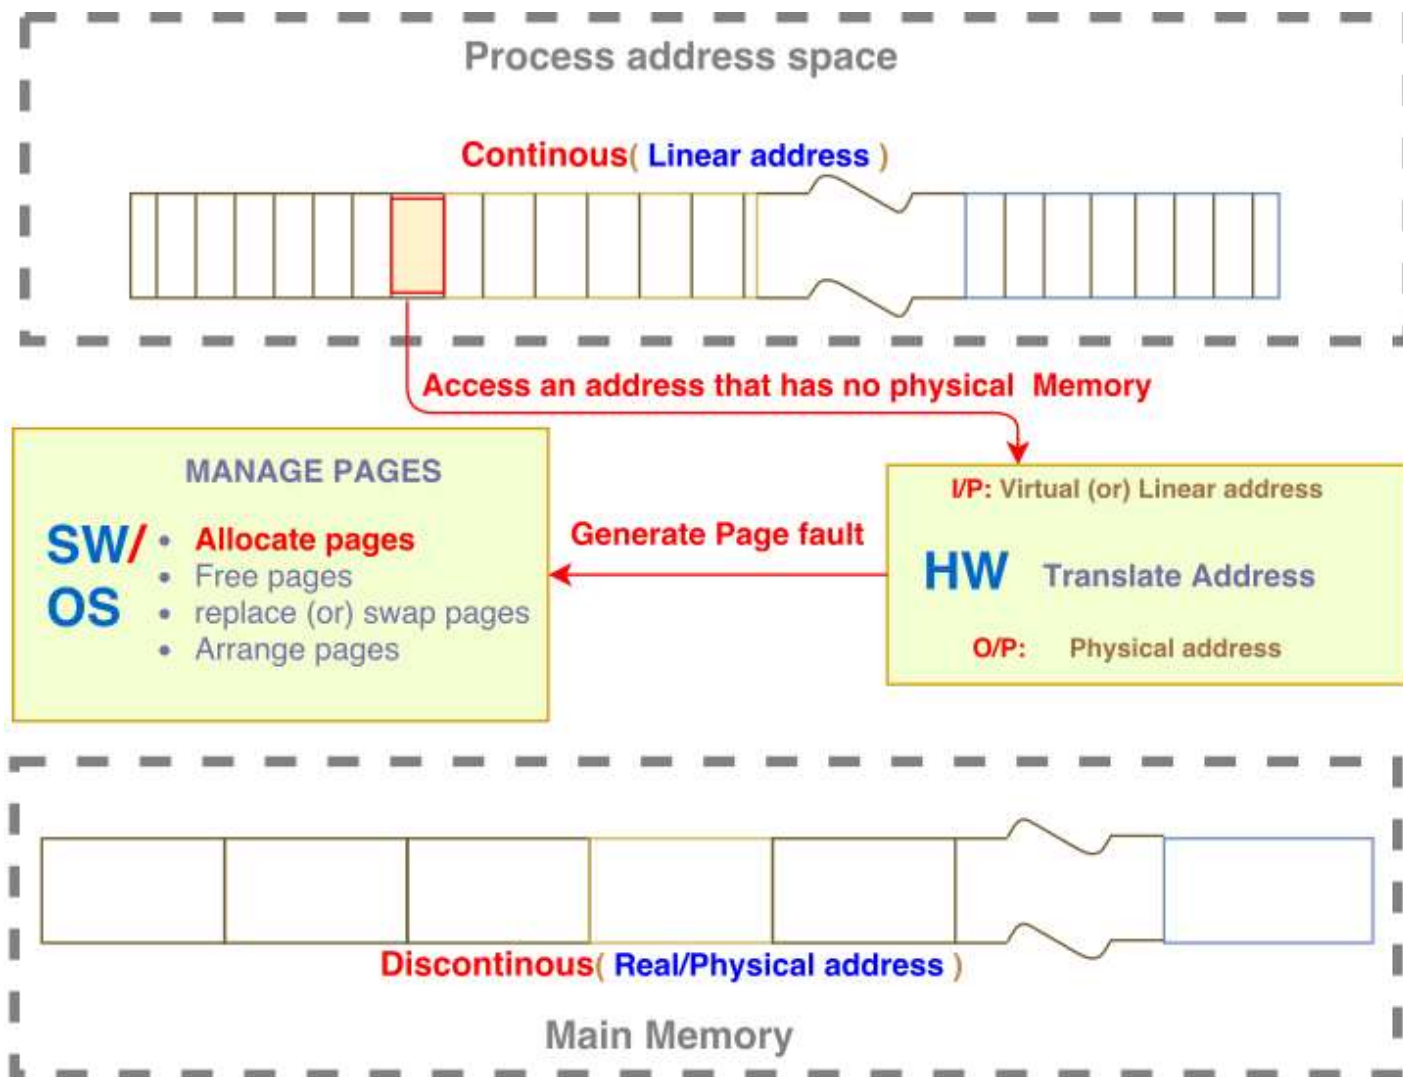
Fetch Policy



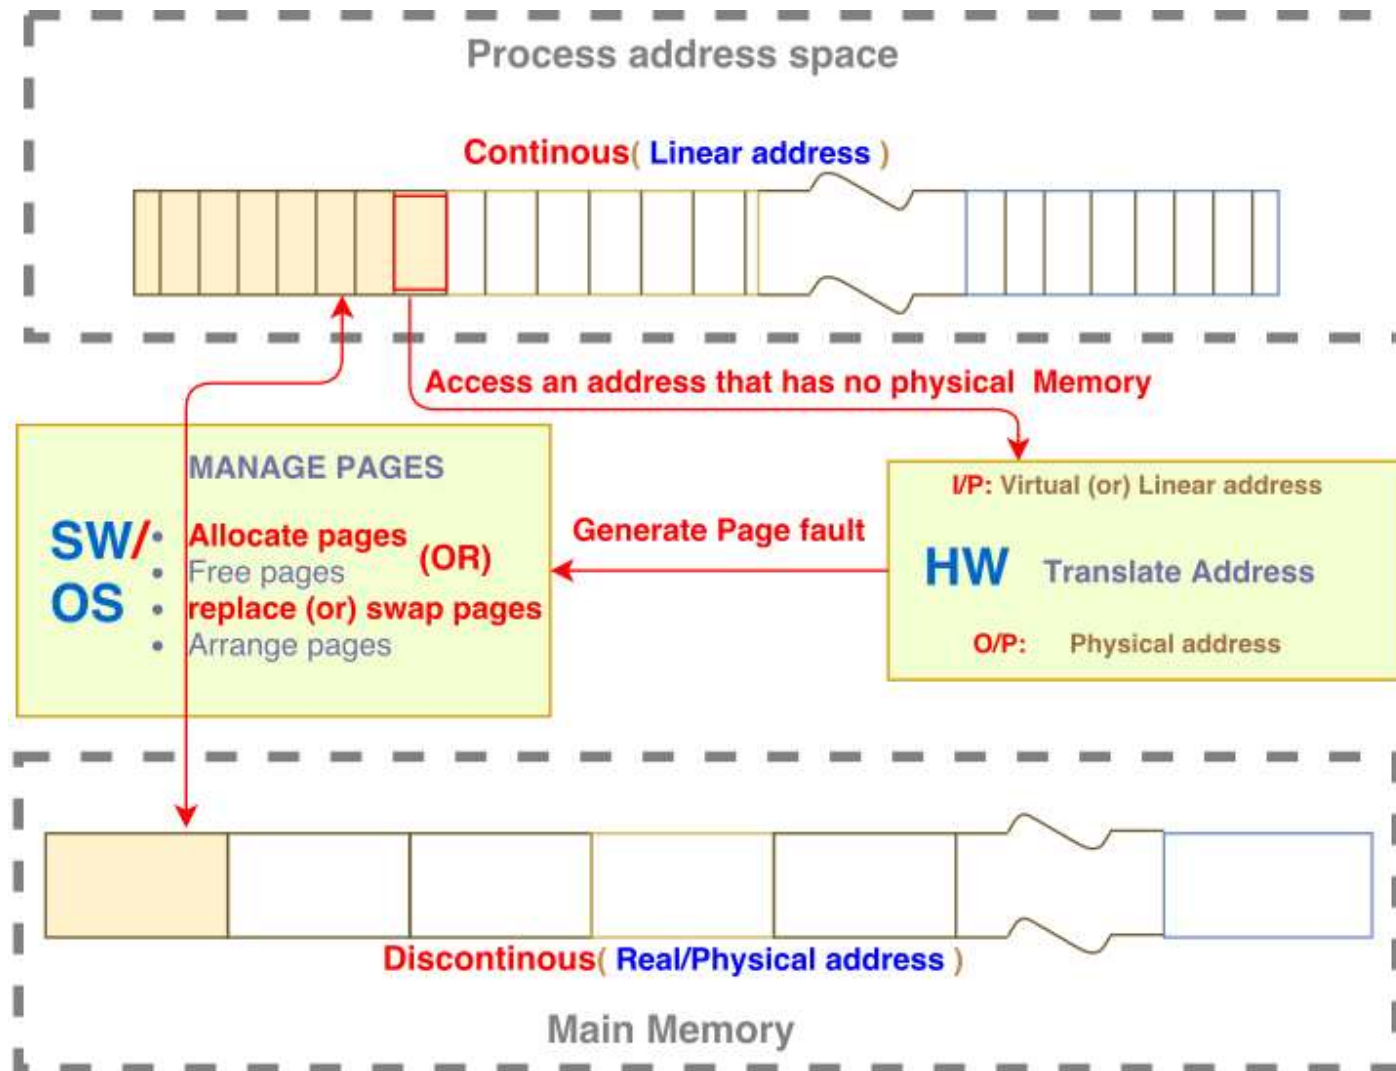
Fetch Policy

- When to fetch the page from auxiliary memory to main memory.
- **2 Approaches:**
 - **Demand Paging** only brings pages into main memory when a **reference** is made to a location on the page.
 - **Prepaging** brings in pages whose use is **anticipated**.

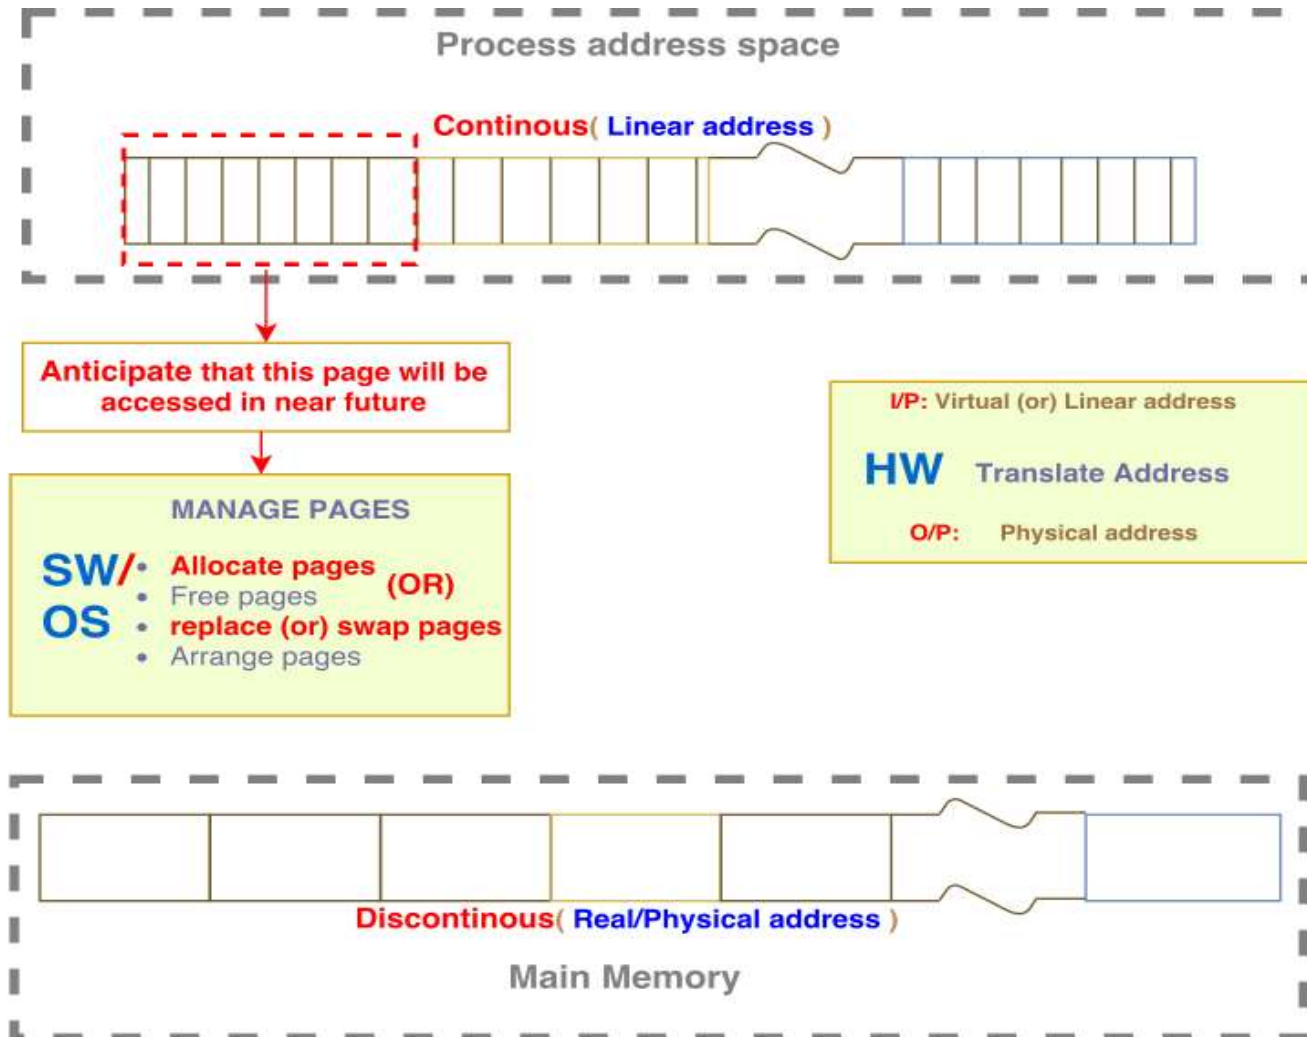
Fetch Policy (Demand Paging)



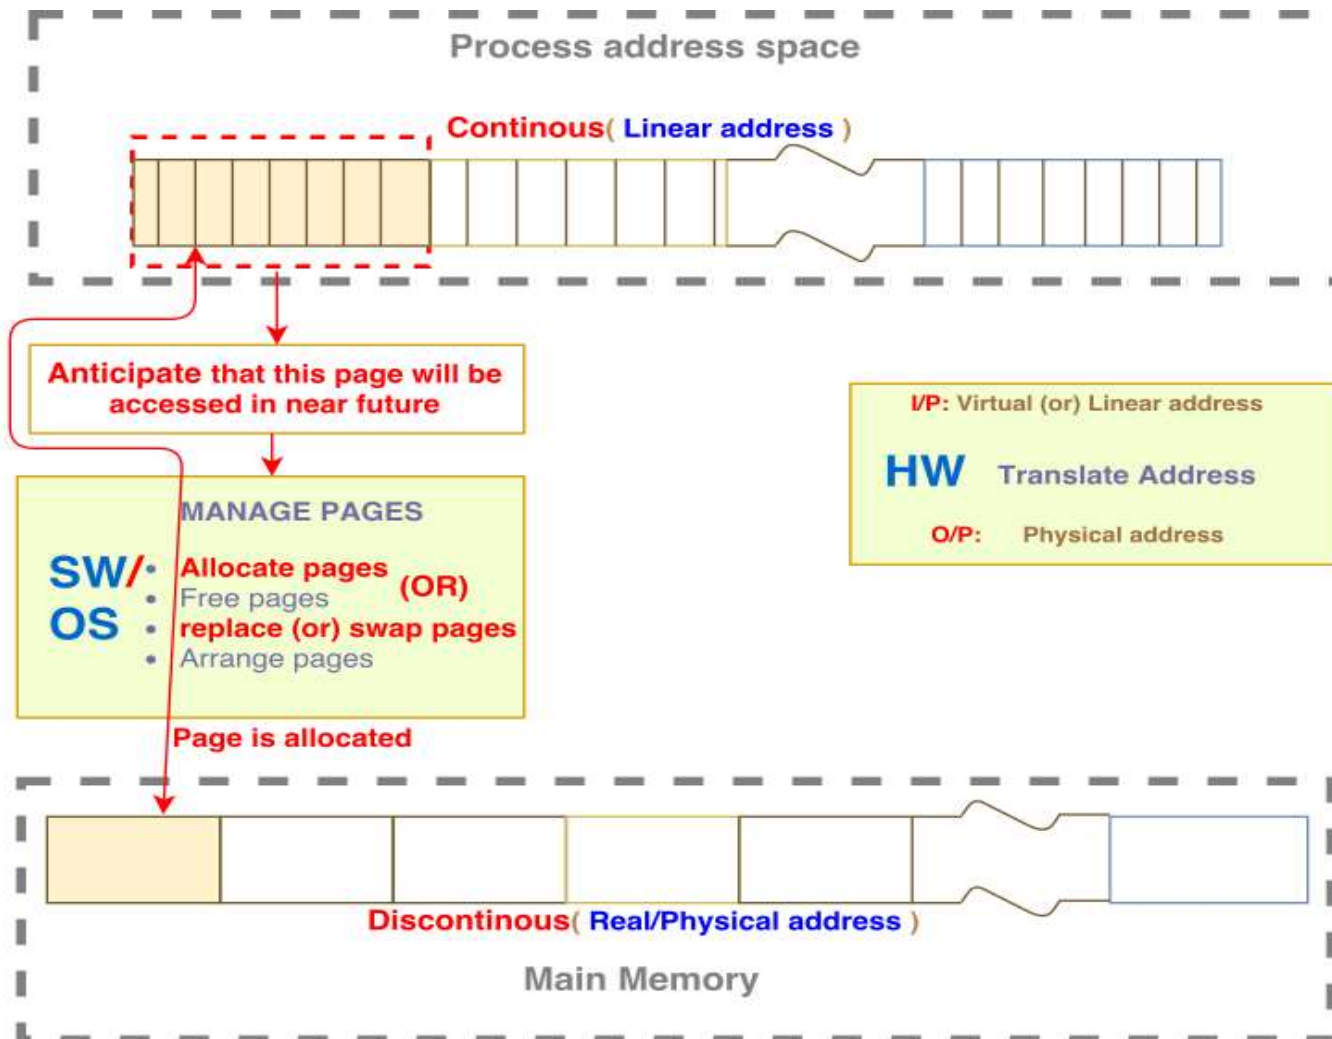
Fetch Policy (Demand Paging)



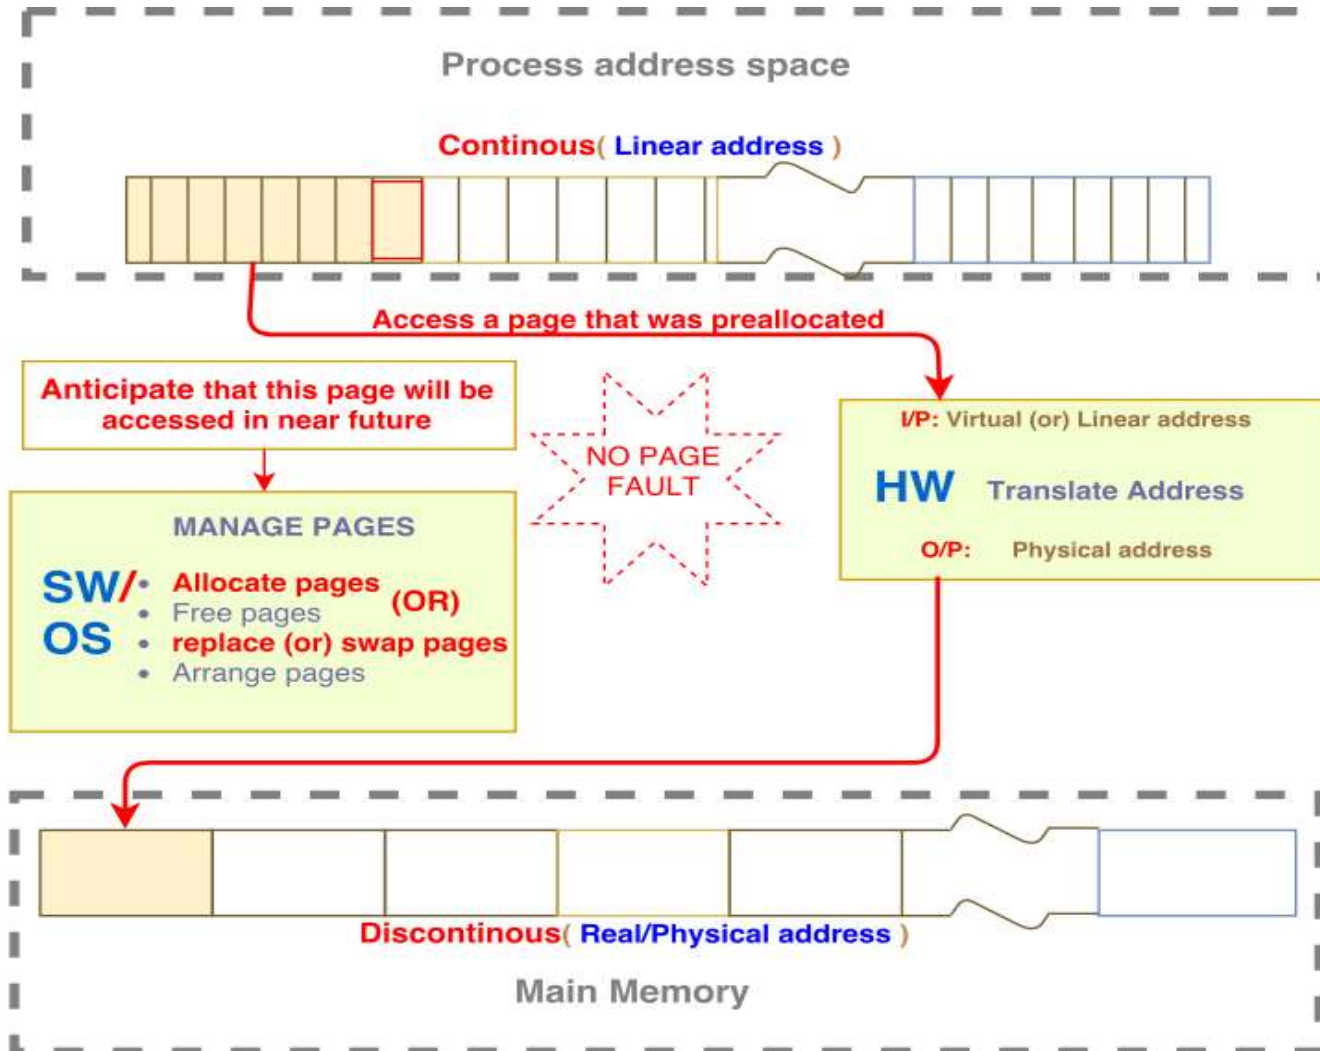
Fetch Policy (Prefetching)



Fetch Policy (Prefetching)



Fetch Policy (Prefetching)



Thank You